

A Global Talent Pool

MATLAB users in education and industry



More job postings require
MATLAB skills than any other
commercial mathematical
computing tool



19 million hits from
search on MATLAB



More Facebook fans than any
other tech computing software
47,000 and growing
(30% increase since Sept '09)

MATLAB Central

User community and
free file exchange –
25,000 downloads/day

40 million annual visits
to mathworks.com

MATLAB Seminar

October 10th, 2012

MathWorks Representatives



Daniel
Armyr

Application
Engineer



Ciamé
Andersson

Education
Account Manager

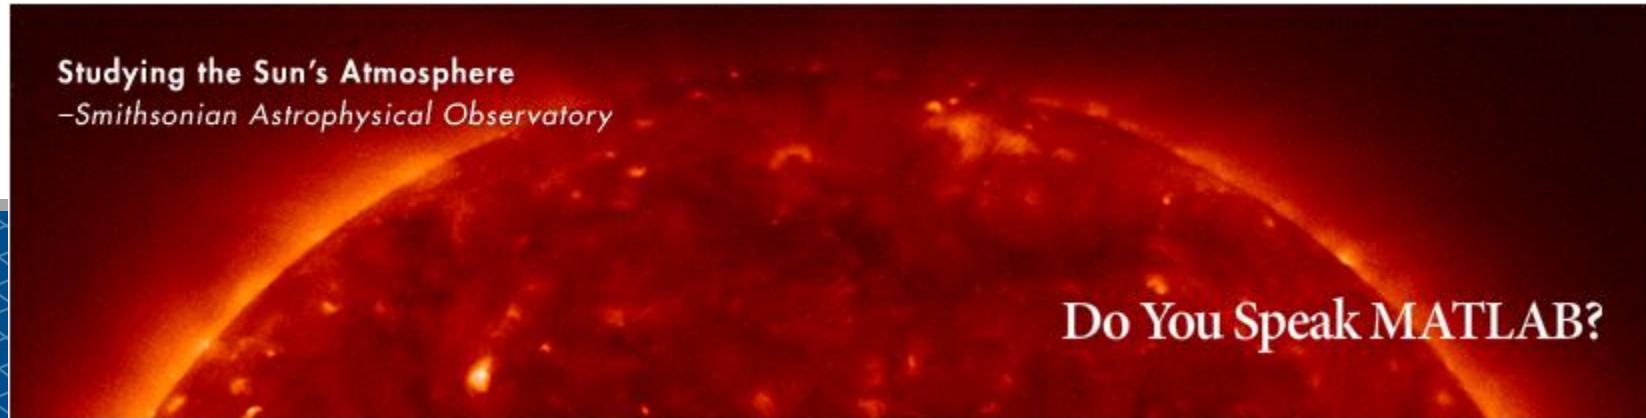
Agenda

- 09:00 - 09:15 Welcome and Introduction
- 09:15 - 10:15 What is MATLAB?
- 10:15 - 10:30 Break
- 10:30 - 11:00 Data Acquisition
- 11:00 - 11:30 Symbolic Computing with MATLAB
- 11:30 - 12:30 Break for lunch
- 12:30 - 13:30 Programming Techniques to Speed up MATLAB
- 13:30 - 13:45 Break
- 13:45 - 14:45 Parallel Computing with MATLAB
- 14:45 - 15:00 Summary and Q&A

What is MATLAB?

or

MATLAB for Excel Users



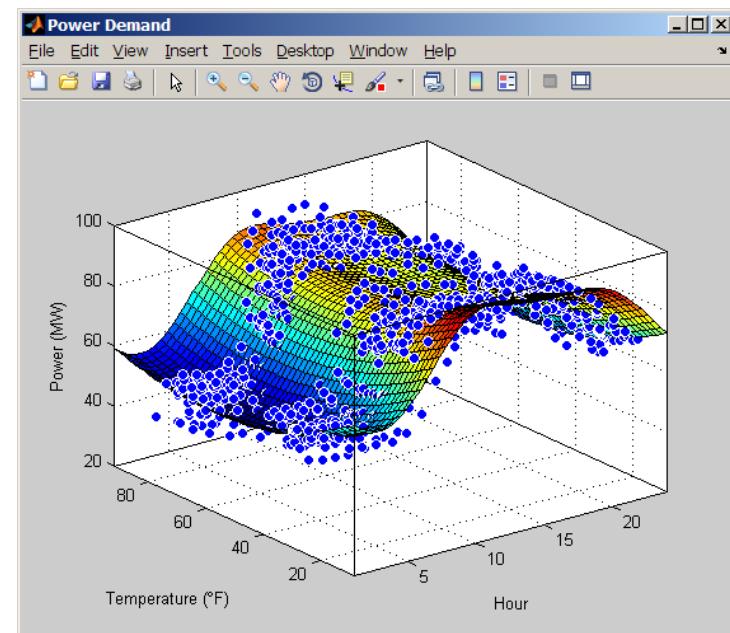
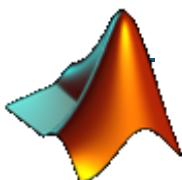
Agenda

-  Demo presentation
 - Data analysis overview
 - Data analysis with MATLAB® and Excel®
 - Break

Demo: Estimate Electricity Usage

Introduction to MATLAB

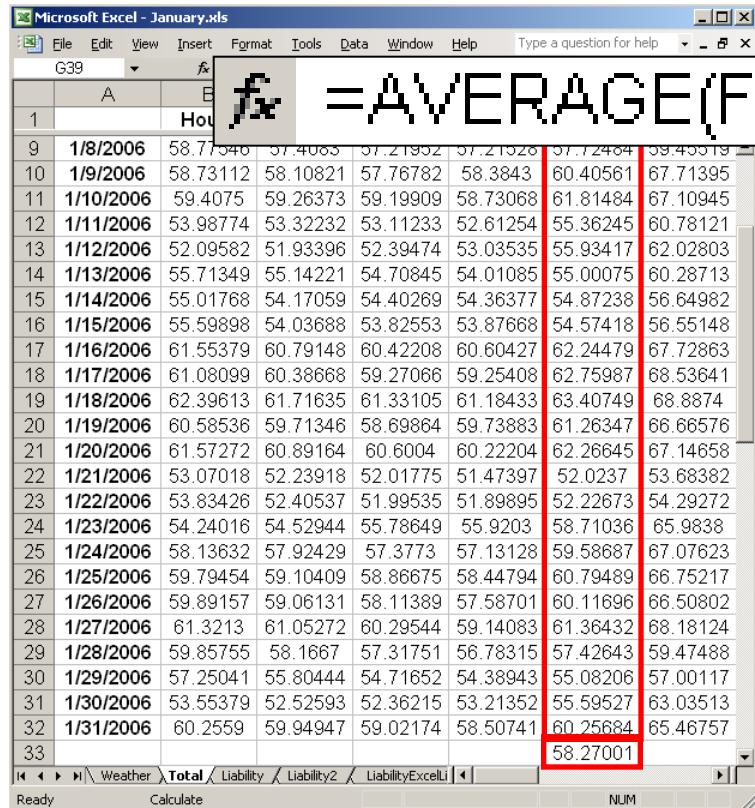
- Goal:
 - Estimate power usage based on time of day and ambient temperature
- Approach:
 - Process historical data of hourly power usage
 - Examine daily, weekly, and monthly trends
 - Develop power usage model
 - Document in a report
 - Link analysis to Excel workbook



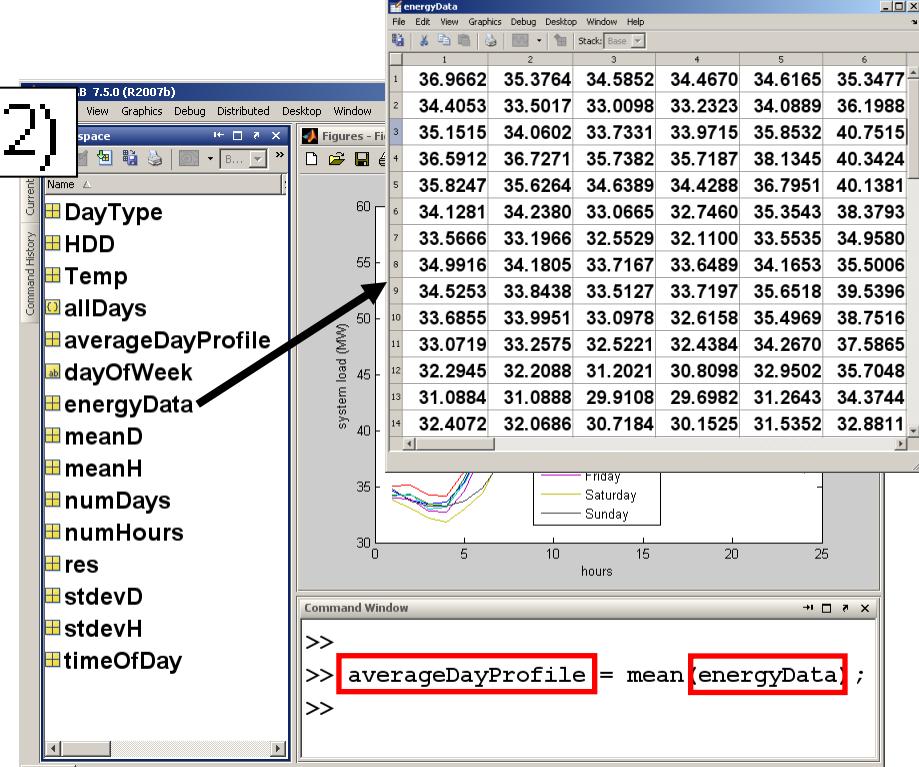
Agenda

- Demo presentation
-  Data analysis overview
- Data analysis with MATLAB® and Excel®
- Break

Approaches with Excel and MATLAB



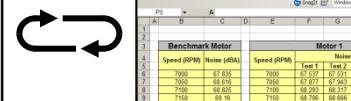
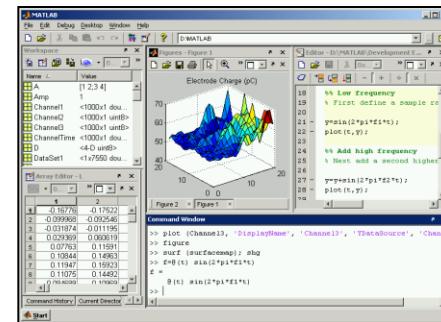
Excel



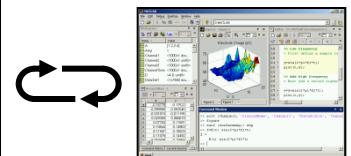
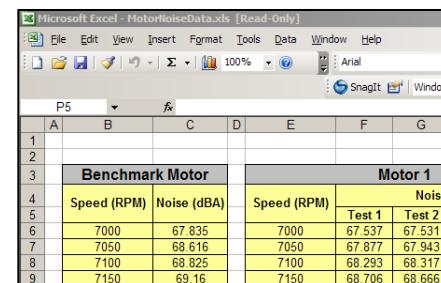
MATLAB

Using MATLAB with Excel

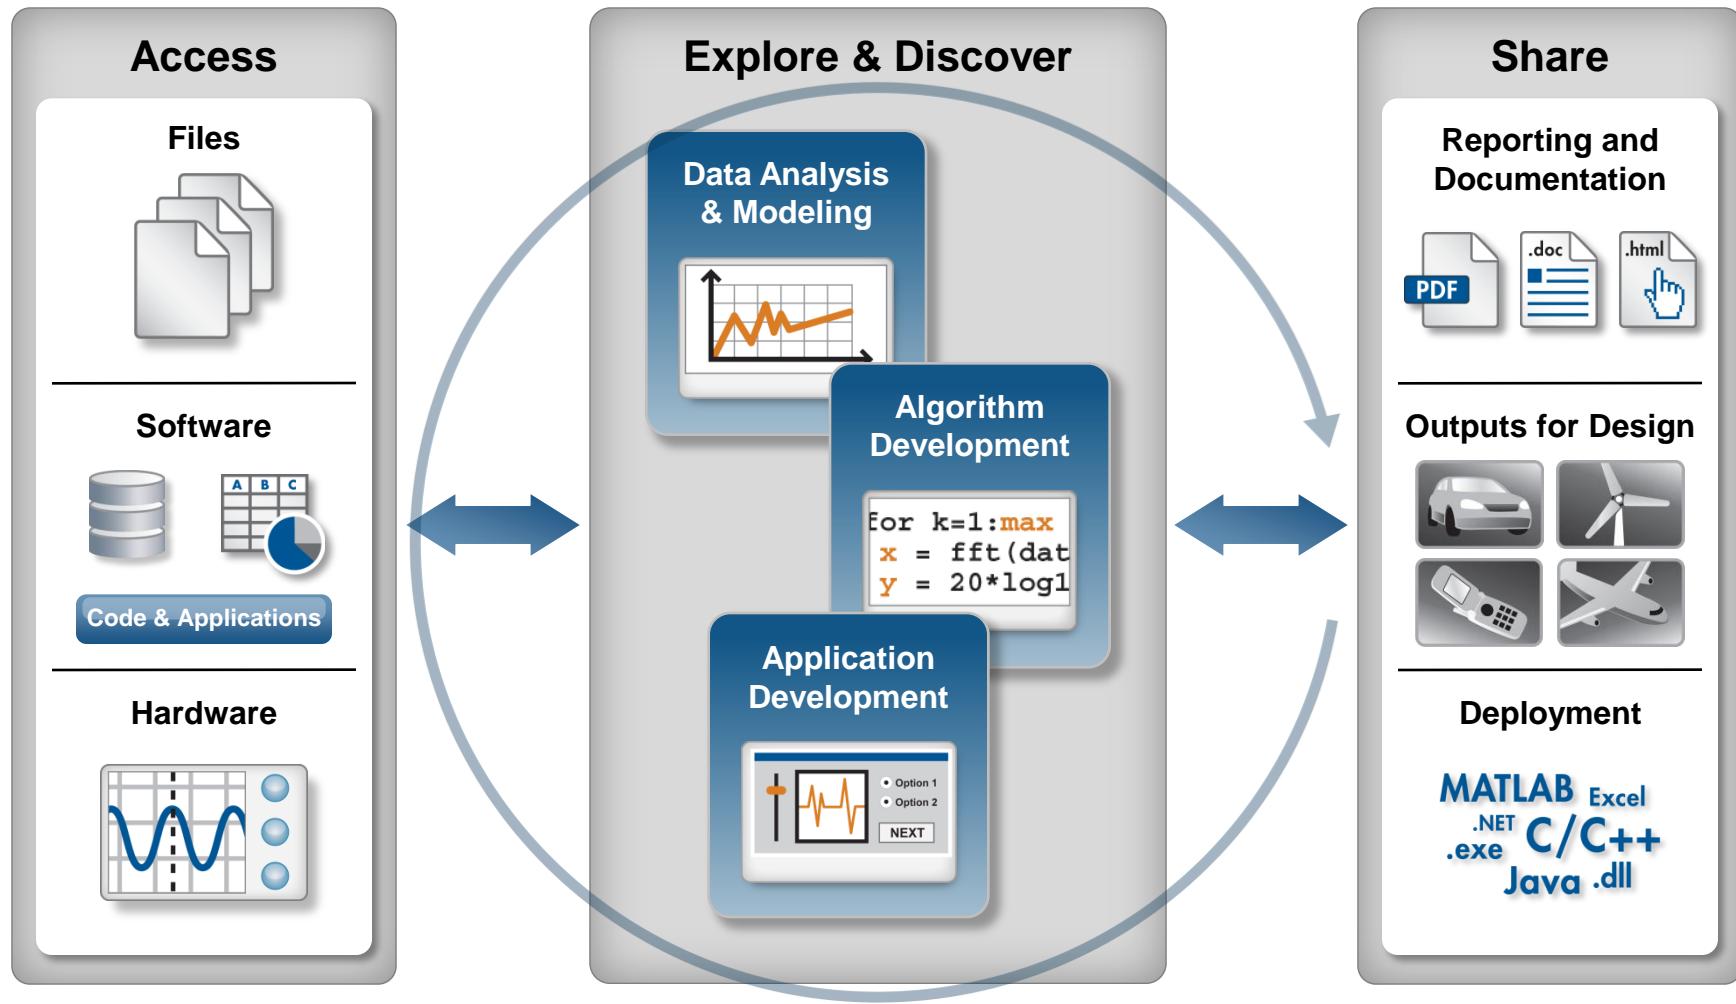
- From MATLAB
 - Read and write Excel data



- From Excel
 - Call MATLAB commands from a cell or Visual Basic script
 - Exchange data with MATLAB
 - Get MATLAB figures

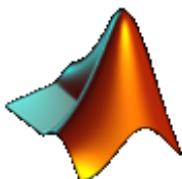


Data Analysis Tasks



Agenda

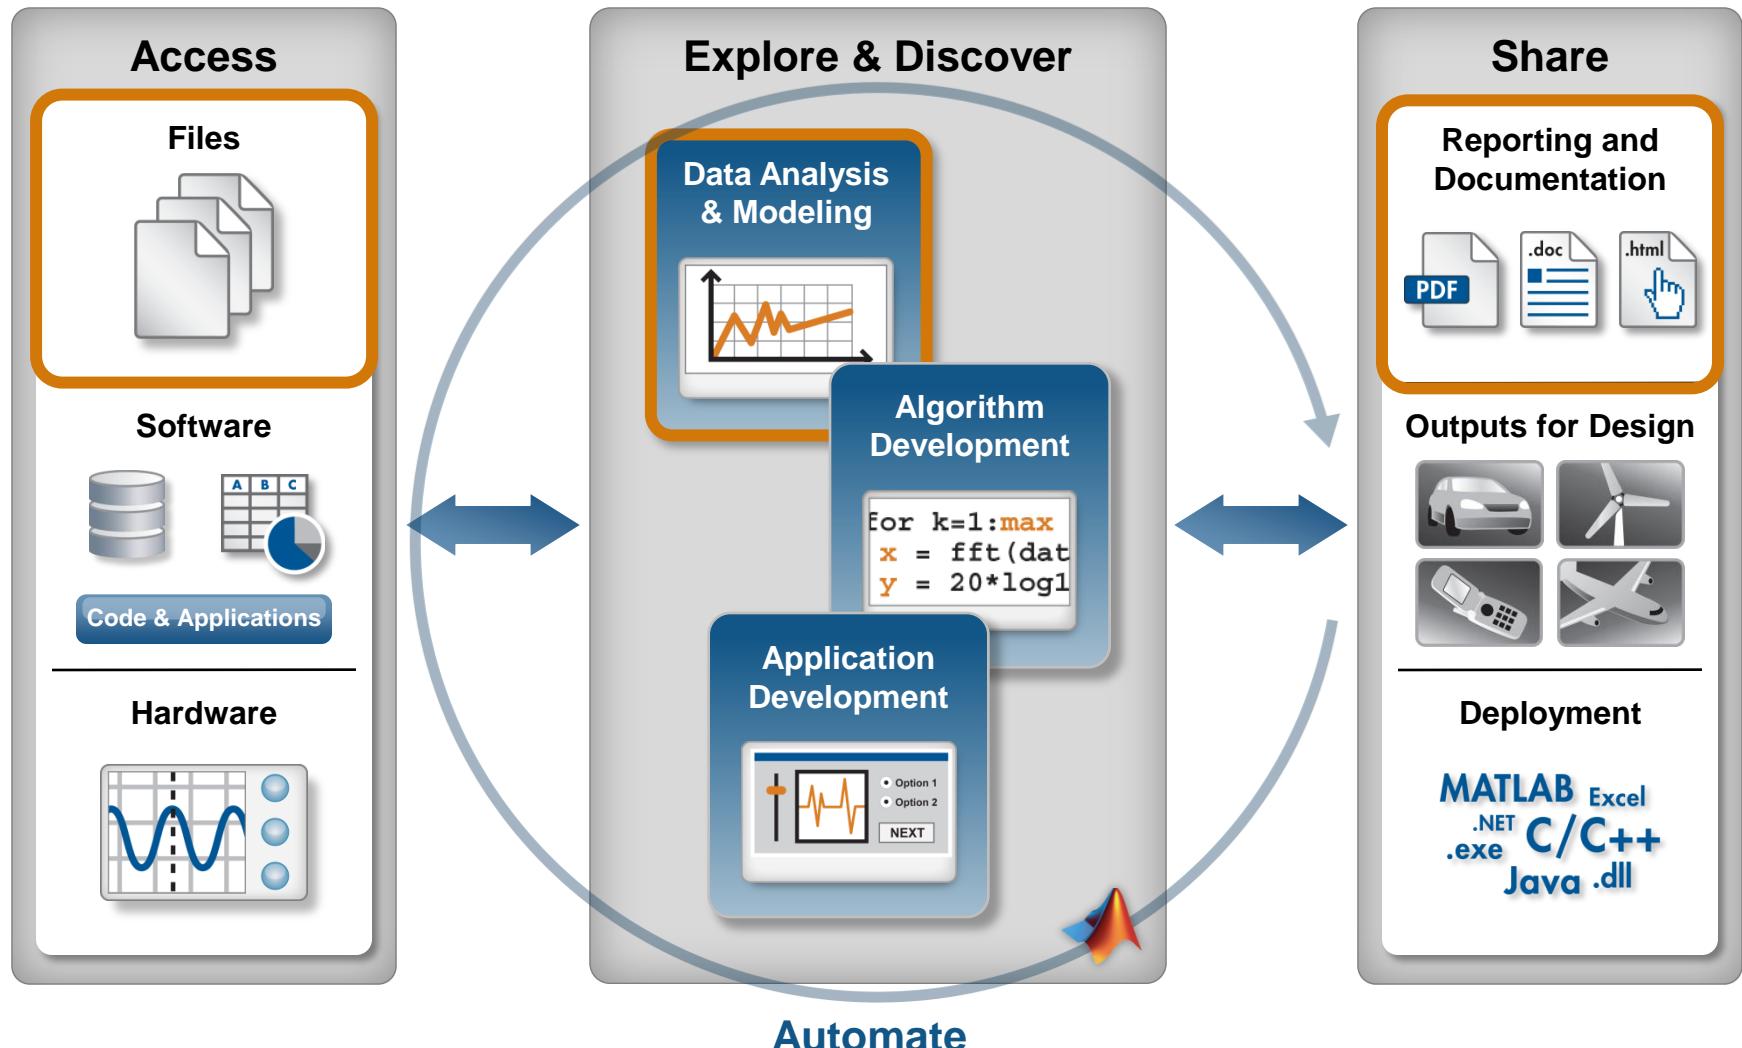
- Demo presentation
- Data analysis overview
-  Data analysis with MATLAB® and Excel®
- Break



Demo: Estimate Electricity Usage

Products Used

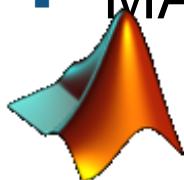
- MATLAB
- Statistics Toolbox
- Curve Fitting Toolbox



Using MATLAB with Excel

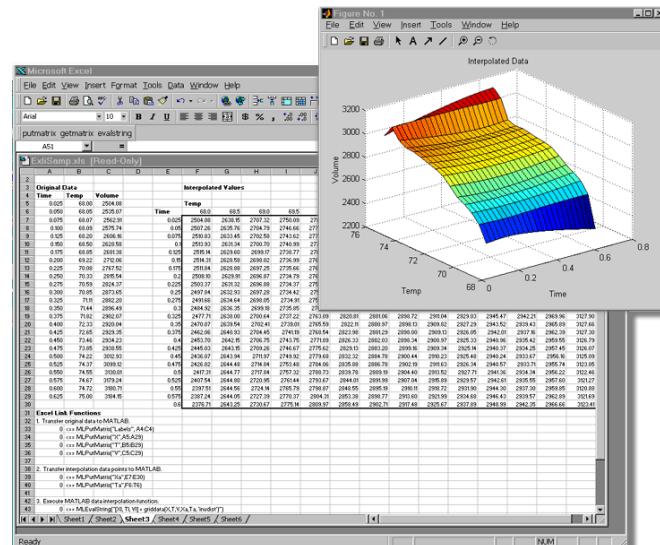
Accessing MATLAB from an Excel Spreadsheet

- **Spreadsheet Link EX**
 - Provides the functions that connect Excel and MATLAB
 - Connection is “live” – requires no intermediate files or inter-process programming
 - **Use Spreadsheet Link EX to:**
 - Pass data and commands to MATLAB
 - Retrieve data and results back into Excel
 - **MATLAB is required**



Products Used

- MATLAB
 - Spreadsheet Link EX

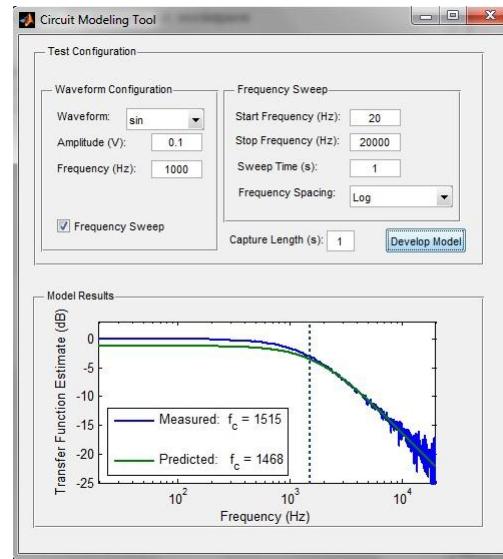
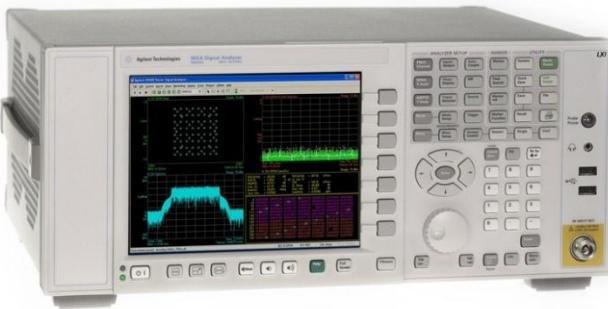


Agenda

- Demo presentation
- Data analysis overview
- Data analysis with MATLAB® and Excel®

 Break

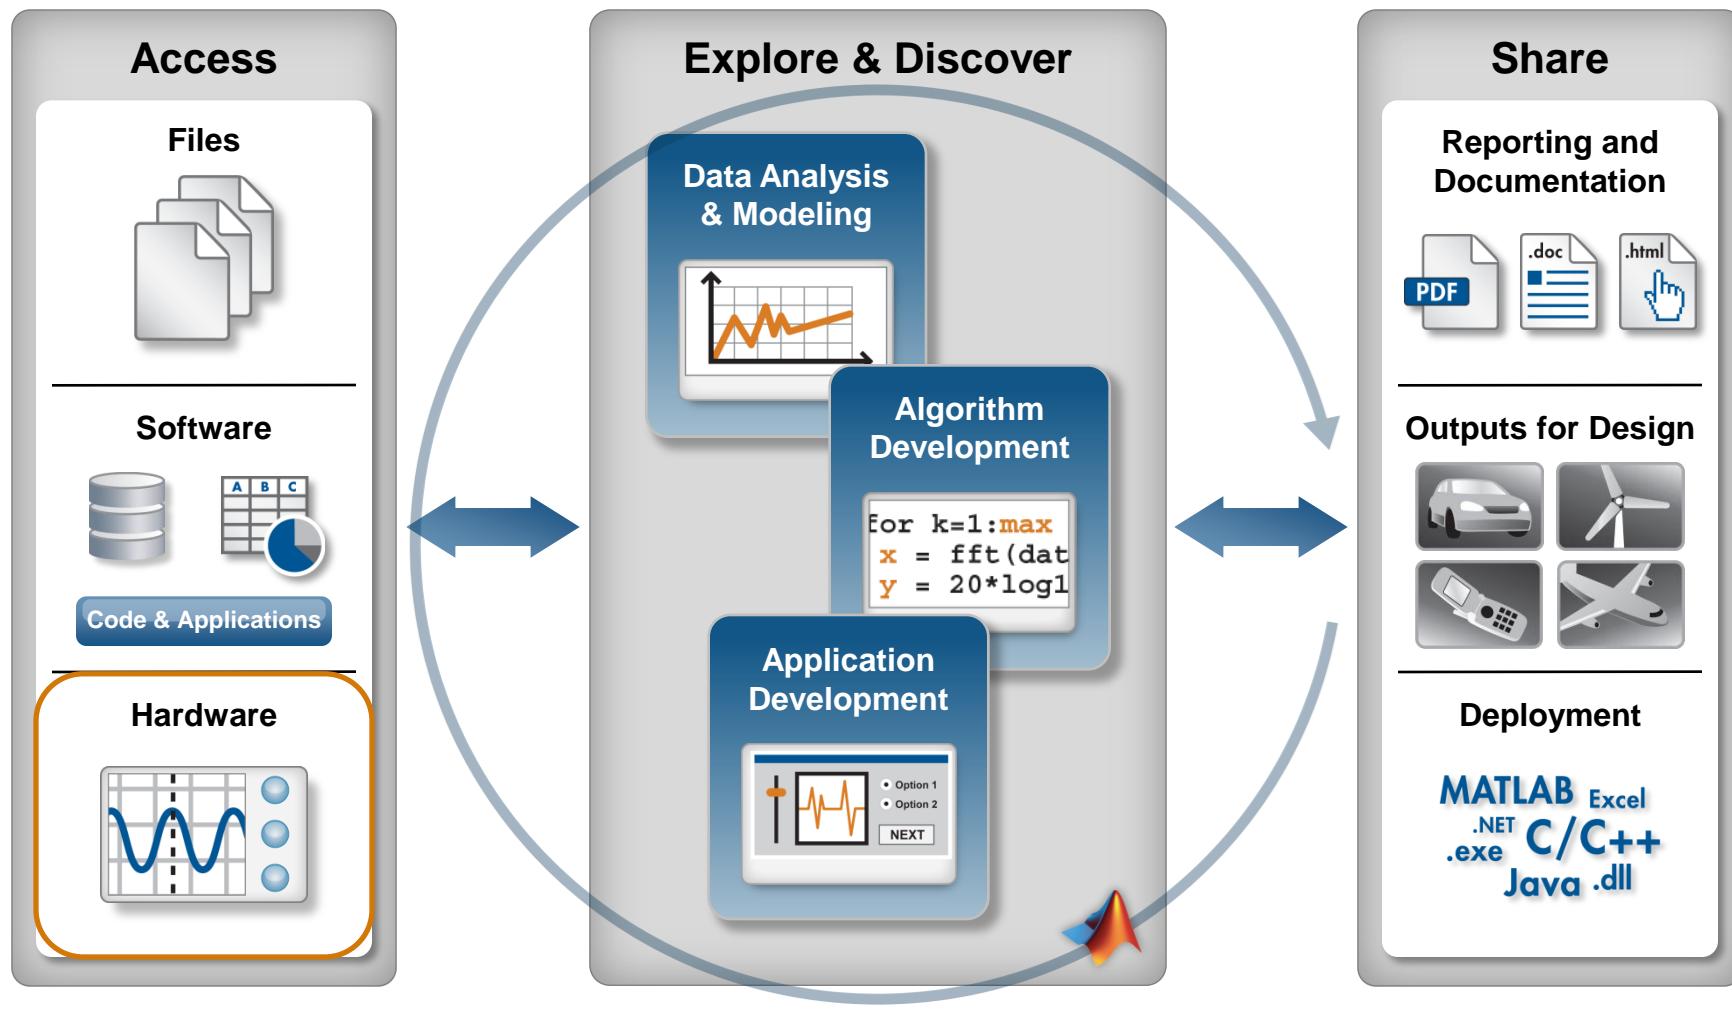
Instrument Control and Data Acquisition with MATLAB



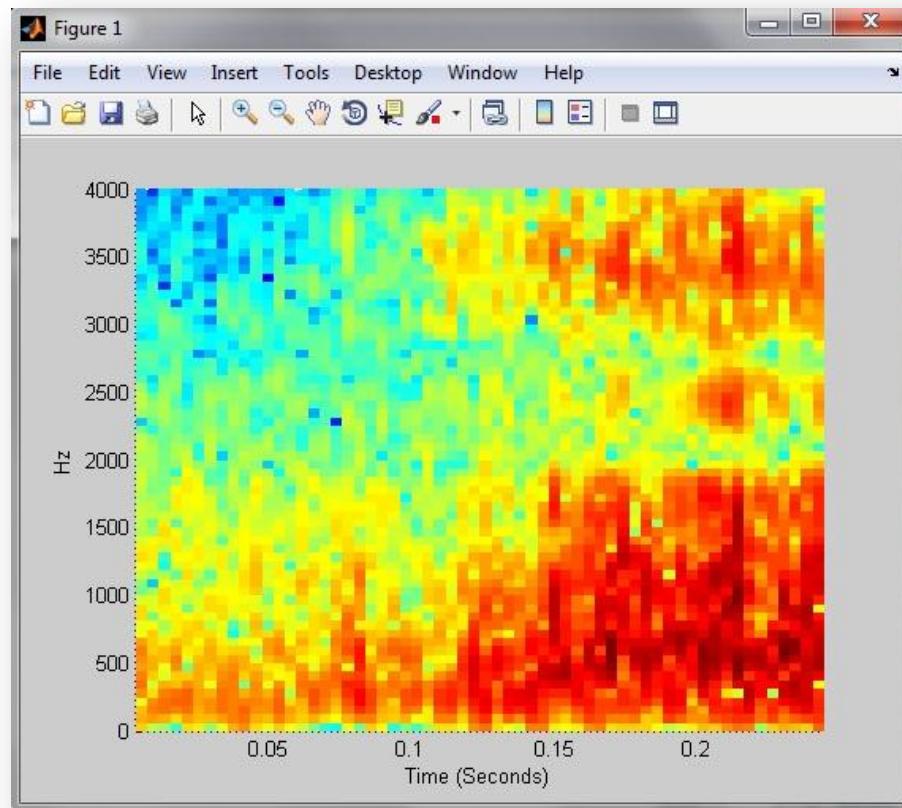
Challenges

1. Separate tools for acquisition and post-test analysis
2. Limited live analysis capabilities
3. Inconsistent support for multiple hardware vendors

Technical Computing Workflow



Demonstration: *Acquiring and Analyzing Data on-the-fly*



MATLAB Connects to Your Hardware Devices

Instrument Control Toolbox

Instruments and RS-232 serial devices



Data Acquisition Toolbox

Plug-in data acquisition devices and sound cards

Image Acquisition Toolbox

Image capture devices



Vehicle Network Toolbox

Vector CAN bus interface devices

OPC Toolbox

OPC DA and HDA servers



Data Acquisition Toolbox:

Supported Hardware

- Agilent
 - National Instruments
 - Advantech
 - CONTEC
 - Data Translation
 - g.tec
 - IOTech
 - Keithley
 - Measurement Computing (MCC)
 - ADLINK
 - Ono Sokki
 - United Electronic Industries
 - VXI Technology
 - Open API
-
- Any PC compatible sound card
 - Parallel Port
-

For a full support listing, visit: www.mathworks.com/products/daq/supportedio.html

Instrument Control Toolbox:

Communication Protocols

GPIB Boards

Agilent Technologies
CONTEC
Keithley
National Instruments

Capital Equipment Corporation (CEC)
IOTech
Measurement Computing
ICS Electronics

VISA Interface (includes Serial, GPIB, VXI, GPIB-VXI, TCP/IP, USB)

Agilent
Rohde & Schwarz

National Instruments
Tektronix

Network Protocols (TCP/IP and UDP)

Serial Port (core functionality in MATLAB)

RS-232, RS-422, RS-485
Expanded capability with Instrument Control Toolbox

For a full listing, visit: www.mathworks.com/products/instrument/hardware/

Image Acquisition Toolbox:

Supported Hardware and Communication Protocols

Analog and Camera Link

BitFlow
Imperx
National Instruments

Data Translation
Matrox Imaging
Teledyne DALSA

DCAM (IIDC 1394)

GenTL

GigE Vision

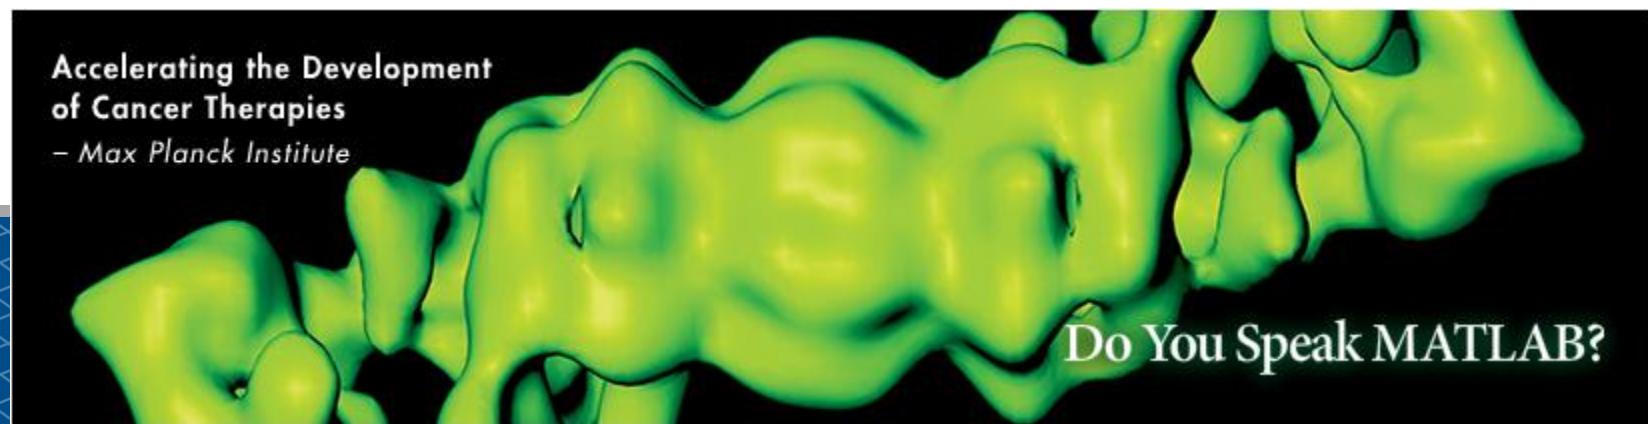
Generic Interface (Webcams)

Video 4 Linux 2
Video for Windows

QuickTime

For a full listing, visit: www.mathworks.com/products/instrument/hardware/

MATLAB at Uppsala University



Uppsala University – License Setup

- TAH Student 
- TAH Campus 
- MDCS

License Access and Administration

<http://www.uu.se/goto/matlab>

<http://www.uu.se/goto/studentmatlab>

Mikael Österberg



mikael.osterberg@angstrom.uu.se

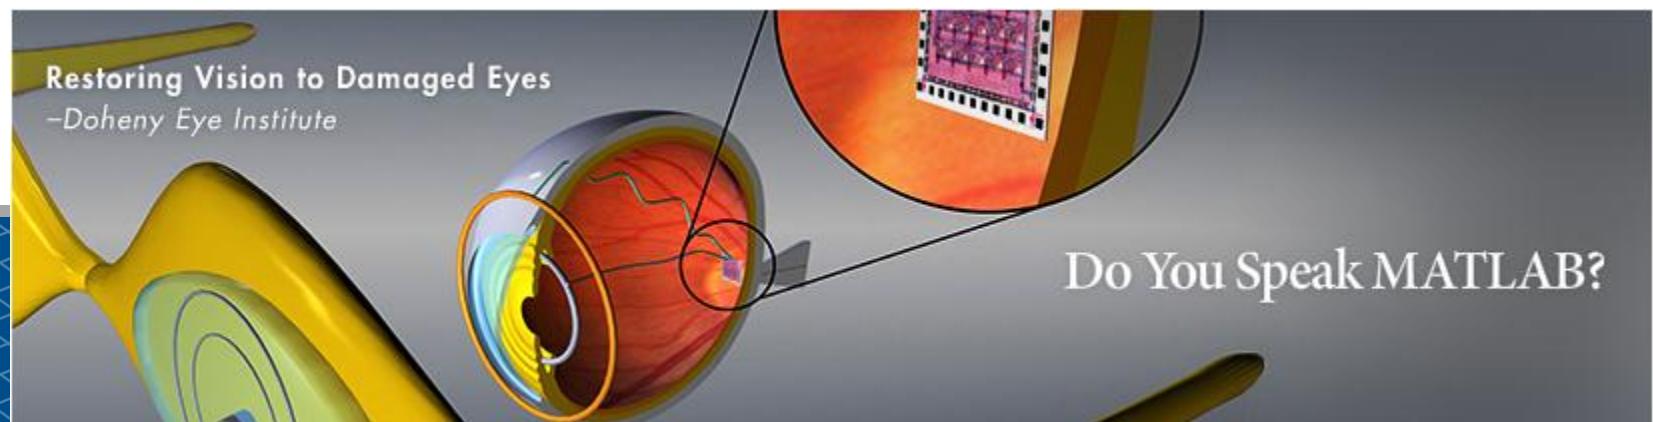


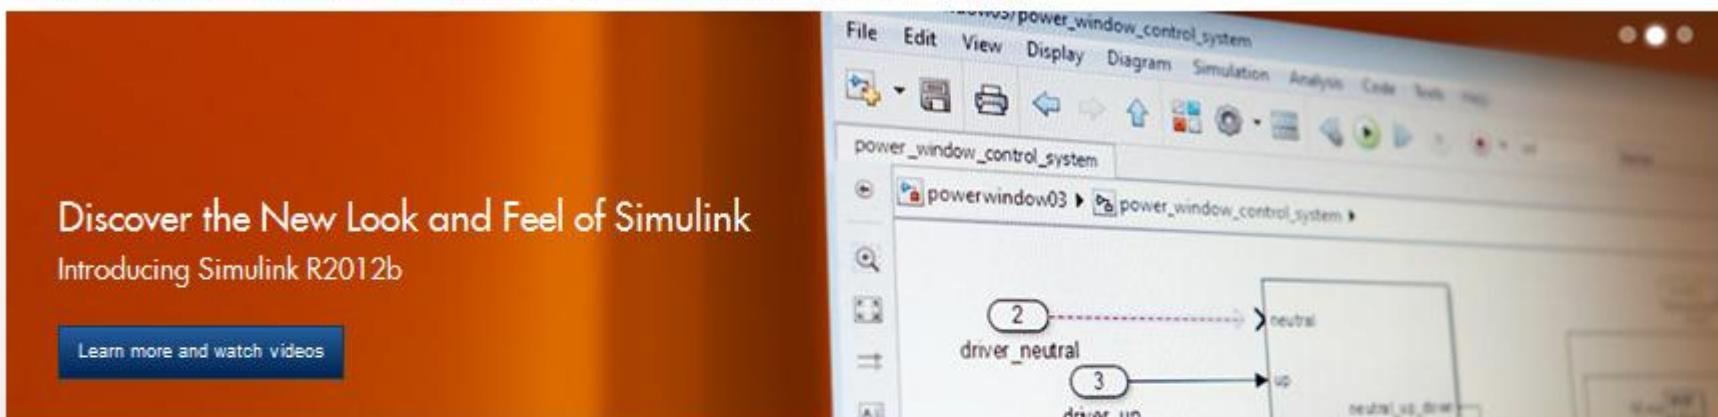
018-471 7989

MathWorks Technical Support & Customer Service

- Tel: 08-5051-6900
 - Monday-Friday
08:30 - 17:00
- www.mathworks.com

Learning more about MATLAB





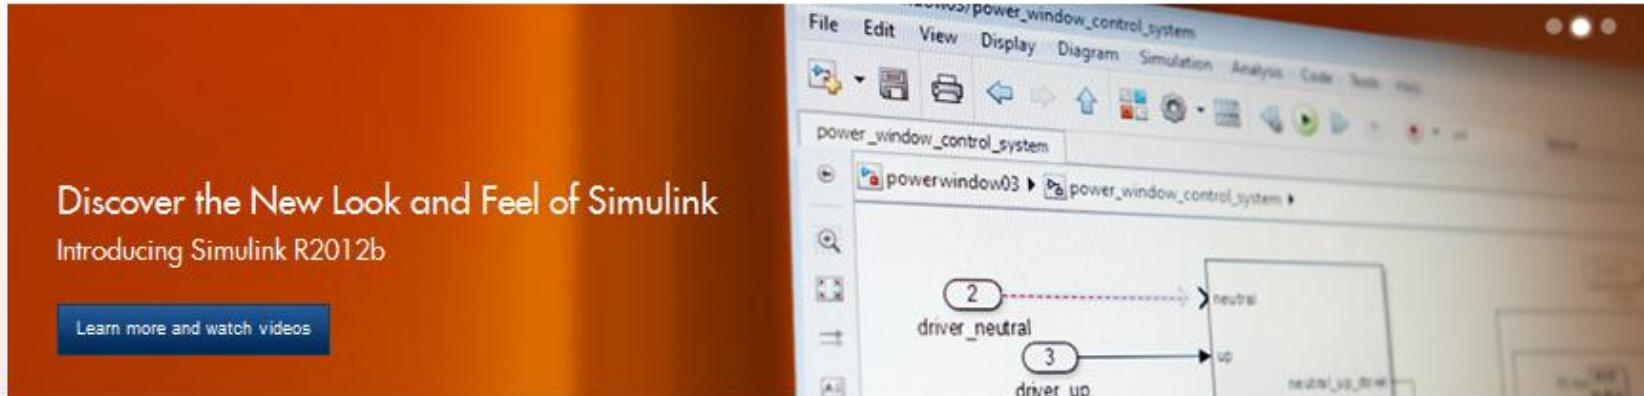
Products  Events Training

Live	On Demand
9 okt Model-Based Design for Hybrid Electric Powertrain... (seminar)	Introduction to Simulink for Control Design
17 okt Computational Finance and Application Deployment with... (seminar)	Model-Based Design for Automation Systems
17 okt Verifying Floating-Point IP Cores on FPGAs with MATLAB... (webinar)	Modeling and Simulating Phased Array Signal Processing Systems
23 okt 4G LTE System Design with MATLAB (webinar)	MATLAB on the Web for the Classroom
24 okt Integrating PK/PD and Mechanistic Modeling with... (webinar)	Acquiring Data from Sensors and Instruments Using MATLAB



MathWorks
ENERGY AND UTILITIES VIRTUAL CONFERENCE
[» Watch videos](#)

Explore all events: webinars, seminars, tradeshows, and conferences



Products Events  Training

Online

03 okt	Stateflow for Logic Driven System Modeling
10 okt	Simulink for System and Algorithm Modeling
23 okt	Polyspace for Code Verification
30 okt	Signal Processing with MATLAB
05 nov	MATLAB Fundamentals

Classroom

Stockholm (City)	MATLAB Fundamentals
Stockholm (City)	MATLAB Programming Techniques
Stockholm (City)	MATLAB Based Optimization Techniques
Stockholm (City)	Simulink for System and Algorithm Modeling
Stockholm (City)	MATLAB Fundamentals



Oct 22nd

New Self-Paced Training
MATLAB Fundamentals
» See course details



Training Course Catalog
» Download now



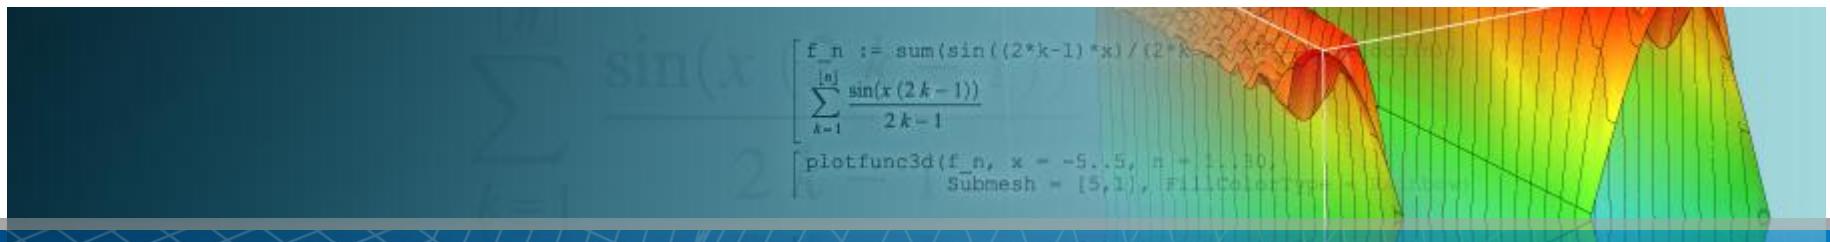
Explore all local classes, e-learning classes, or other training options



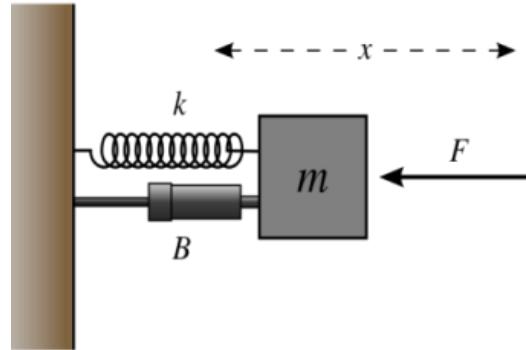
Ciamé Andersson
Education Account Manager

ciame.andersson@mathworks.com 
08-505 169 61 

Symbolic Computing with MATLAB



Overview of Symbolic Computing



Governing equation:

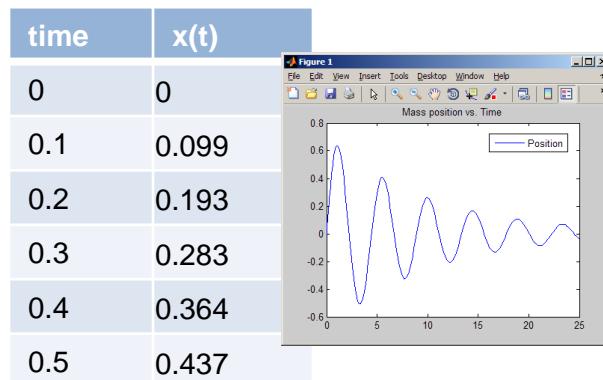
$$F = mx''(t) + Bx'(t) + kx(t)$$

Initial conditions:

$$x(0) = 0 \quad x'(0) = 0$$

Numeric:

Approximate solution in vector form



Symbolic:

Exact solution in form of analytical expression

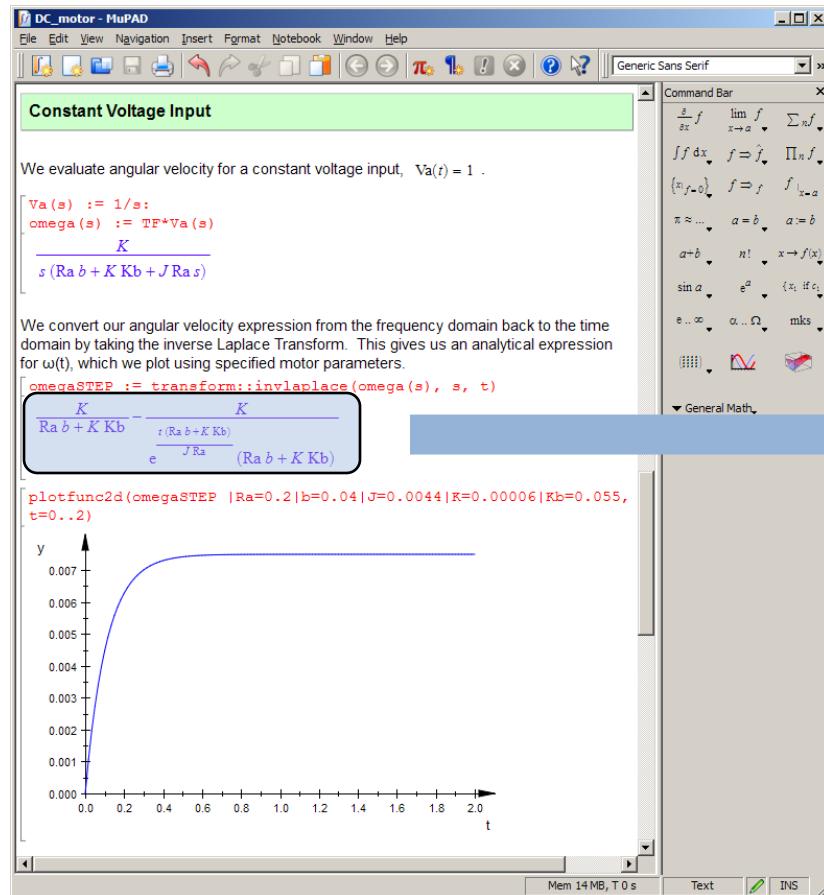
$$x(t) = \frac{m}{e^{\frac{t(B-\sqrt{B^2-4km})}{2m}}} - \frac{m}{e^{\frac{t(B+\sqrt{B^2-4km})}{2m}}}$$

The terms $\sqrt{B^2 - 4km}$ are circled in red.

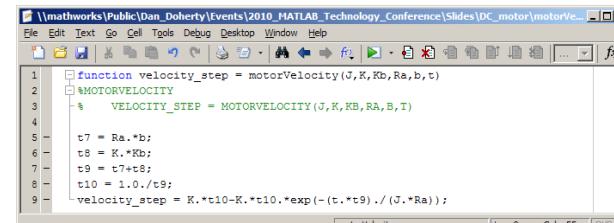
Benefits

1. Insight into how parameters affect solution
2. Often more efficient than numeric implementation

Integrating Symbolic Results with MATLAB and Simulink



with
MATLAB

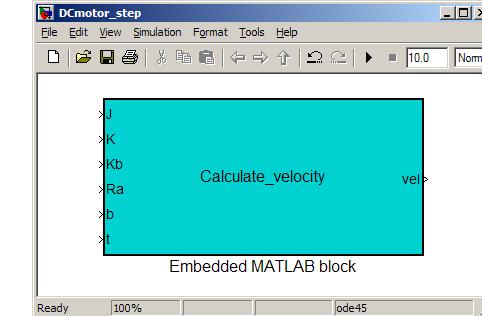


```

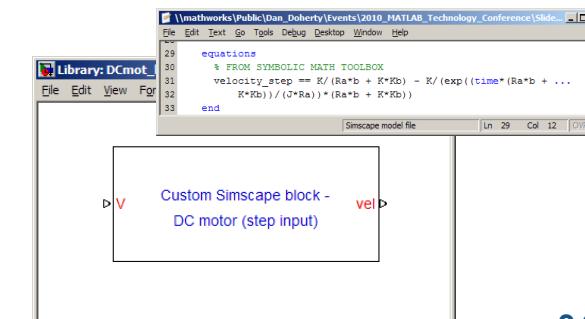
1 function velocity_step = motorVelocity(J,K,Kb,Ra,b,t)
2 %MOTORVELOCITY
3 % VELOCITY_STEP = MOTORVELOCITY(J,K,Kb,Ra,B,T)
4
5 t7 = Ra.*b;
6 t8 = K.*Kb;
7 t9 = t7+t8;
8 t10 = 1.0./t9;
9 velocity_step = K.*t10.*t10.*exp(-(t.*t9)./(J.*Ra));

```

with
Simulink



with
Simscape



```

29
30 equations
31 % FROM SYMBOLIC MATH TOOLBOX
32 velocity_step == K/(Ra*b + K*Kb) - K/(exp((time*(Ra*b + ...
33 K*Kb))/(J*Ra))*(Ra*b + K*Kb))
end

```

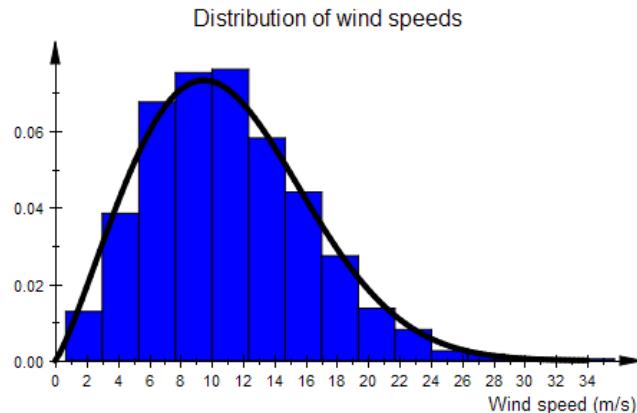
Custom Simscape block - DC motor (step input)

Product Demonstration

Modeling power generated by a wind turbine

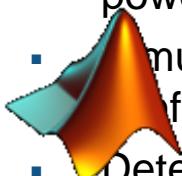
Objective

- Determine wind turbine design that produces optimal power at a given site



Problem-solving steps

- Derive analytical model for average power produced by wind turbine
- Simulate model in MATLAB for different configurations
- Determine wind turbine that optimizes power at given site



Evaluating the Average Power Delivered by a Wind Turbine

Background



The total power delivered to a wind turbine can be estimated by taking the derivative of the wind's kinetic energy. This results in the following expression:

$$(1) \quad P_w = \frac{\rho \cdot A \cdot u^3}{2}$$

where A = swept area of turbine blades, in m^2
 ρ = air density, in kg/m^3
 u = wind speed, in m/s

The process of converting wind power to electrical power results in efficiency losses, as described in the diagram below.

```

graph LR
    Pw[Power Input] --> Turbine[Turbine  
Cp]
    Turbine --> Transmission[Transmission  
Ct]
    Transmission --> Generator[Generator  
Cg]
    Generator --> Pe[Output Power]
  
```

The electrical power output of a practical wind turbine can be described using the following equation:

Summary

- Develop algorithms and models with symbolic math
 - Gain efficiency
 - Gain insight through analytical solutions
 - Perform interactive calculation and documentation via the notebook interface
 - Integrate analytical results in MATLAB

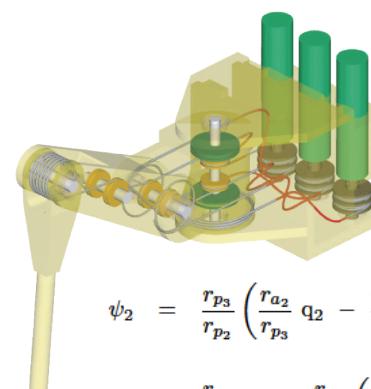


$$C = S\Phi(d_1) - Ke^{-rT}\Phi(d_2)$$

where

$$d_1 = \frac{\ln(S/K) + (r + \sigma^2/2)T}{\sigma\sqrt{T}}$$

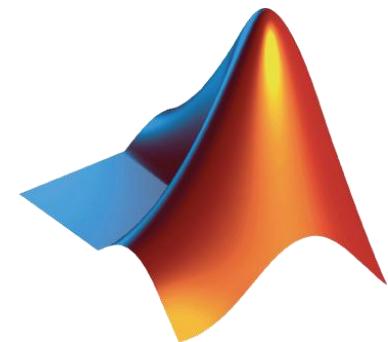
$$d_2 = \frac{\ln(S/K) + (r - \sigma^2/2)T}{\sigma\sqrt{T}} = d_1 - \sigma\sqrt{T}.$$



$$\psi_2 = \frac{r_{p_3}}{r_{p_2}} \left(\frac{r_{a_2}}{r_{p_3}} q_2 - \frac{r_{a_1}}{r_{p_1}} q_1 \right)$$

$$\phi_2 = \frac{r_{a_2}}{r_{p_2}} q_2 + \frac{r_{a_1}}{r_{p_1}} \left(\frac{r_{p_2} - r_{p_3}}{r_{p_2}} \right) q_1$$

Speeding up MATLAB Applications



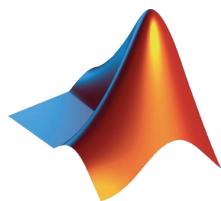
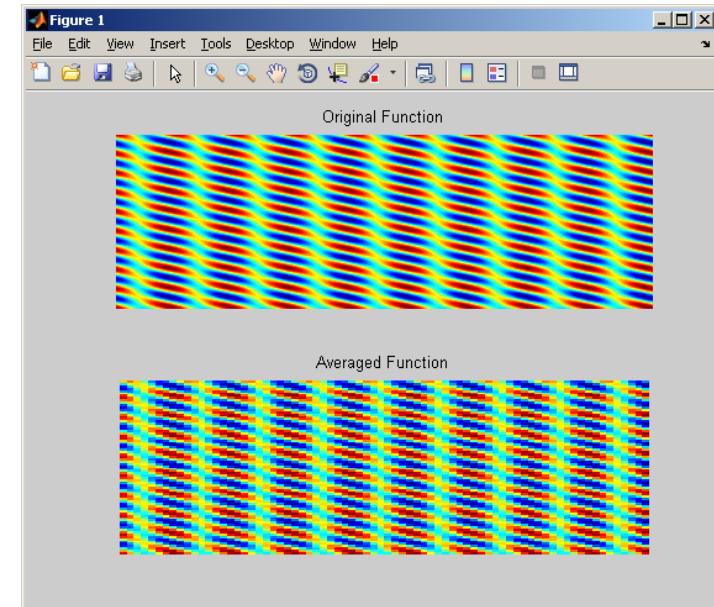
Agenda



- Leveraging the power of vector and matrix operations in MATLAB®**
 - Demonstration: Preallocation and vectorization
- **How does MATLAB® store and provide access to its' variables?**
 - Demonstration: Row- vs. Column-major indexing
 - Demonstration: Logical indexing
- **Addressing bottlenecks**
 - Demonstration: Profiling code in MATLAB®
- **Memory considerations**
 - Demonstration: Copy-on-write
- **Summary and Q&A**

Example: Block Processing Images

- Evaluate function at grid points
- Reevaluate function over larger blocks
- Compare the results
- Evaluate code performance



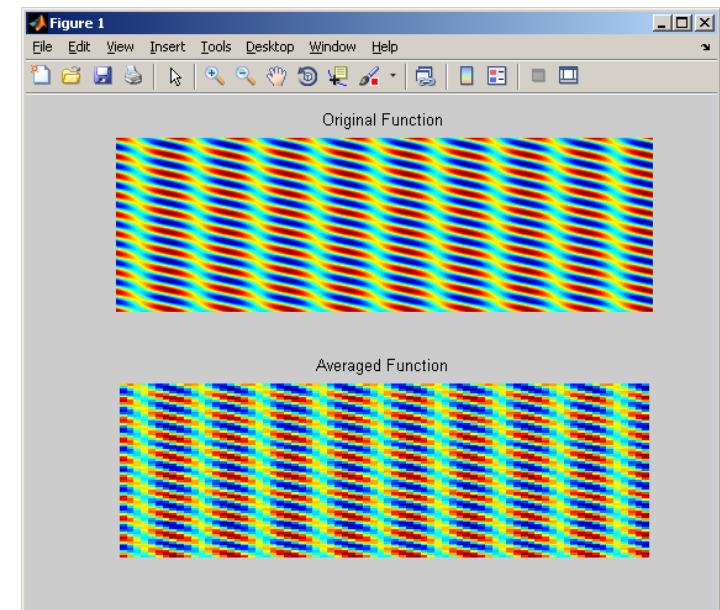
Summary of Example

- Used built-in timing functions

```
>> tic  
>> toc
```

- Used Code Analyzer to find suboptimal code

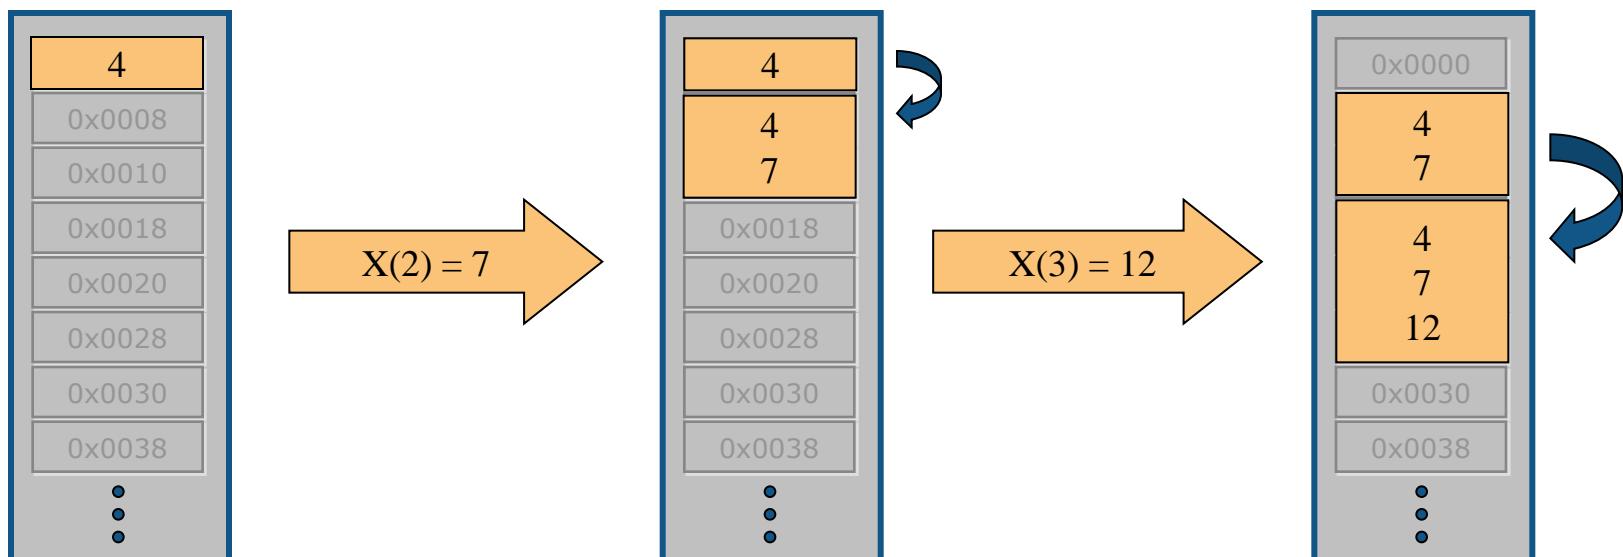
- Preallocated arrays
- Vectorized code



Effect of Not Preallocating Memory

```
>> x = 4  
>> x(2) = 7  
>> x(3) = 12
```

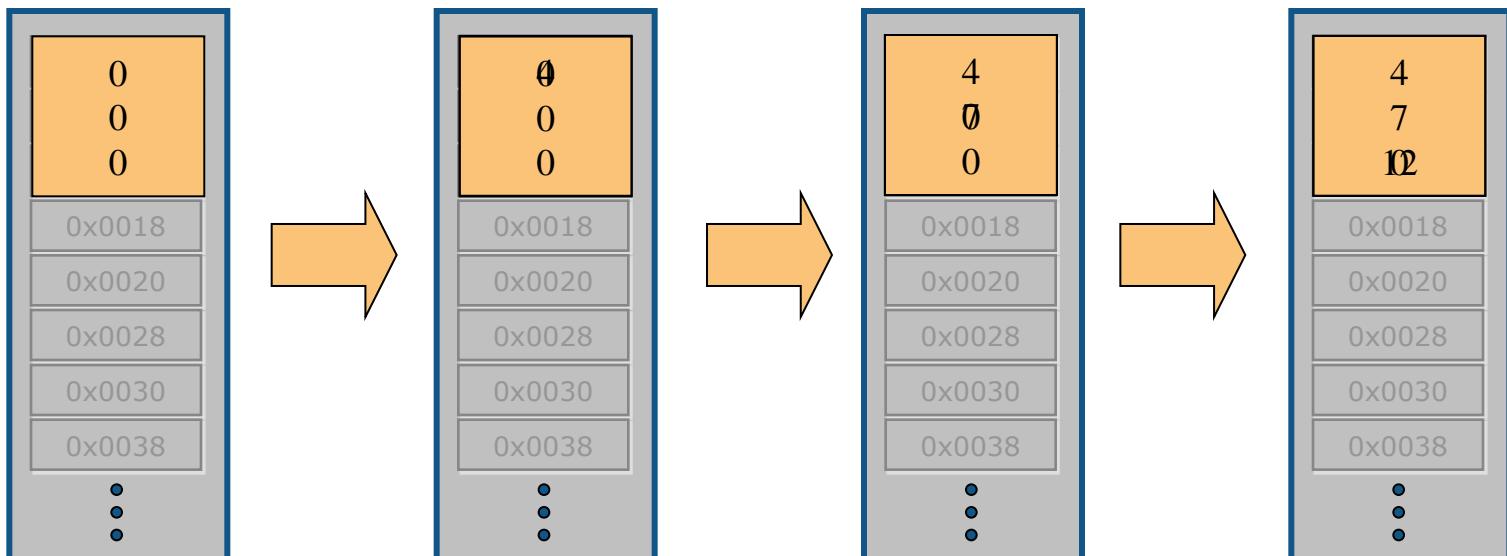
Resizing
Arrays is
Expensive



Benefit of Preallocation

```
>> x = zeros(3,1)  
>> x(1) = 4  
>> x(2) = 7  
>> x(3) = 12
```

Reduced
Memory
Operations



Agenda

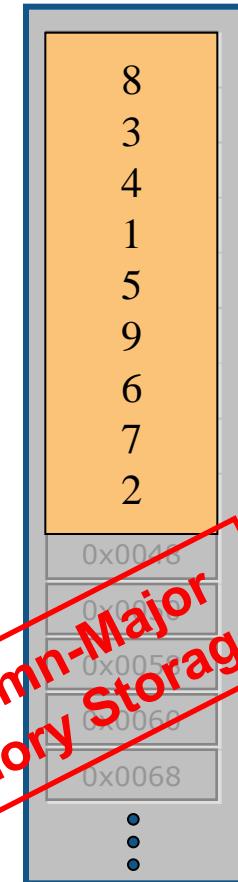
- Leveraging the power of vector and matrix operations in MATLAB®
 - Demonstration: Preallocation and vectorization
-  How does MATLAB® store and provide access to its' variables?
 - Demonstration: Row- vs. Column-major indexing
 - Demonstration: Logical indexing
- Addressing bottlenecks
 - Demonstration: Profiling code in MATLAB®
- Memory considerations
 - Demonstration: Copy-on-write
- Summary and Q&A

Data Storage of MATLAB Arrays

What is the fastest way to process
MATLAB matrices with for loops?
(i.e. de-vectorize)

- a) Down the columns
- b) Along the rows
- c) Doesn't matter

```
>> x = magic(3)  
x =  
     8   1   6  
     3   5   7  
     4   9   2
```



See the June 2007 article in “The MathWorks News and Notes”:

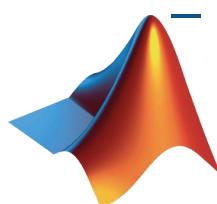
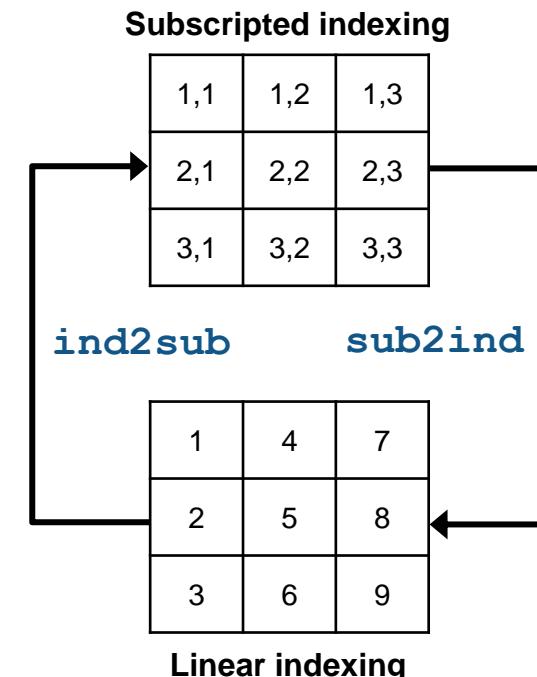
http://www.mathworks.com/company/newsletters/news_notes/june07/patterns.html

Speed and Memory Usage

- Balance vectorization and memory usage
 - Use **bsxfun** instead of functions such as **repmat**
 - Reduce size of arrays to smaller blocks for block processing
- Consider using sparse matrices
 - *Less Memory*: Store only nonzero elements and their indices
 - *Faster*: Eliminate operations on zero elements
 - Blog Post - Creating Sparse Finite Element Matrices
<http://blogs.mathworks.com/loren/2007/03/01/creating-sparse-finite-element-matrices-in-matlab/>

Indexing into MATLAB Arrays

- Subscripted
 - Access elements by rows and columns
- Linear
 - Access elements with a single number
- Logical
 - Access elements with logical operations or mask



MATLAB Underlying Technologies

- Commercial libraries
 - BLAS: Basic Linear Algebra Subroutines (multithreaded)
 - LAPACK: Linear Algebra Package
 - etc.
- JIT/Accelerator
 - Improves looping
 - Generates on-the-fly multithreaded code
 - Continually improving

BLAS and LAPACK
require contiguous
arrays

Other Best Practices

- Minimize dynamically changing path

```
>> addpath(...)  
>> fullfile(...)
```

instead of cd(...)

- Use the functional load syntax

```
>> x = load('myvars.mat')
```

```
x =
```

```
    a: 5  
    b: 'hello'
```

instead of load('myvars.mat')

- Minimize changing variable class

```
>> x = 1;  
>> xnew = 'hello';
```

instead of x = 'hello';

Agenda

- Leveraging the power of vector and matrix operations in MATLAB®
 - Demonstration: Preallocation and vectorization
- How does MATLAB® store and provide access to its' variables?
 - Demonstration: Row- vs. Column-major indexing
 - Demonstration: Logical indexing

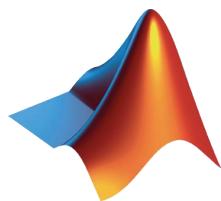
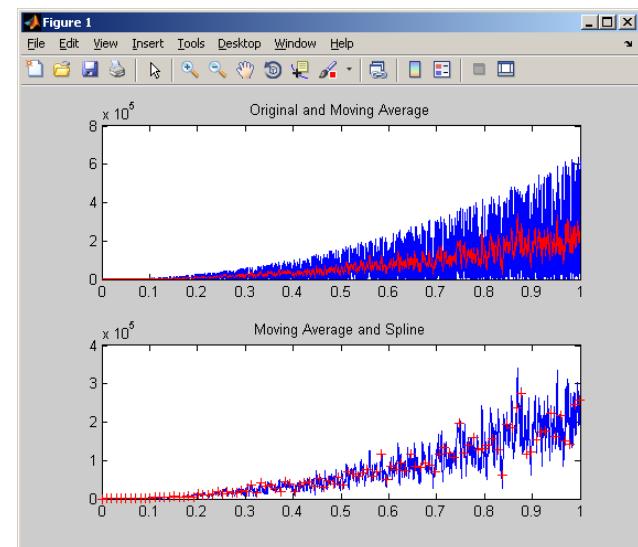


Addressing bottlenecks

- Demonstration: Profiling code in MATLAB®
- Memory considerations
 - Demonstration: Copy-on-write
- Summary and Q&A

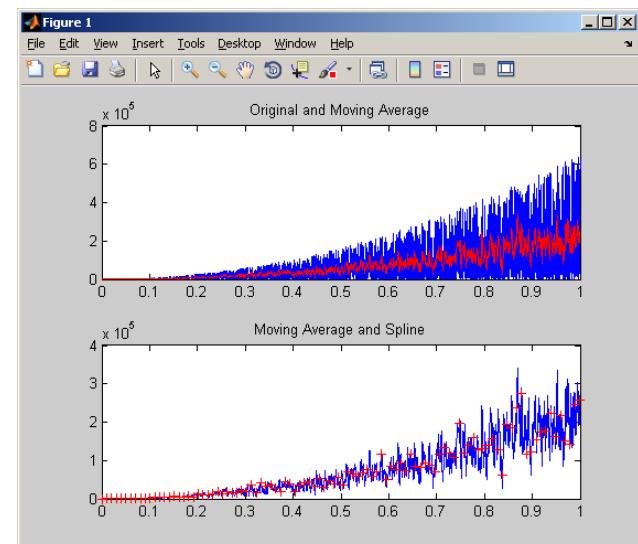
Example: Fitting Data

- Load data from multiple files
- Extract a specific test
- Fit a spline to the data
- Write results to Microsoft Excel



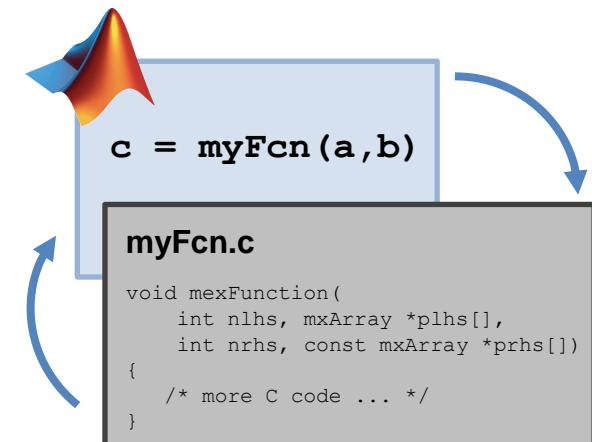
Summary of Example

- Used profiler to analyze code
- Targeted significant bottlenecks
- Reduced file I/O
- Reused figure



Acceleration using MEX (MATLAB Executable)

- Call C or Fortran code directly from MATLAB
 - Integrate existing code using MEX API
 - Auto-generate C-based MEX files from MATLAB code using MATLAB Coder
- Speed-up factor will vary
 - May see speedup for state-based for-loops
 - May not see a speedup when MATLAB code is
 - Using multithreaded computations
 - Using optimized libraries (BLAS, FFTW, etc.)



Agenda

- Leveraging the power of vector and matrix operations in MATLAB®
 - Demonstration: Preallocation and vectorization
- How does MATLAB® store and provide access to its' variables?
 - Demonstration: Row- vs. Column-major indexing
 - Demonstration: Logical indexing
- Addressing bottlenecks
 - Demonstration: Profiling code in MATLAB®
-  Memory considerations
 - Demonstration: Copy-on-write
- Summary and Q&A

MATLAB and Memory

Lazy copying (copy-on-write):

```
C{1} = rand(3000,3000); %72Mb
```

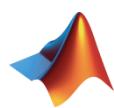
```
C{2} = rand(3000,3000); %72Mb
```



```
C_new = C; %0 Mb
```

```
C{1}(1,1) = 2; %72Mb
```

Arrays are stored
in contiguous
memory



In-place Optimizations

When does MATLAB do calculations “**in-place**”?

```
function y = myfunc(x)
y = sin(2*x.^2+3*x+4);
```

```
function x = myfuncInPlace(x)
x = sin(2*x.^2+3*x+4);
```

```
function inplaceOptimization
x = randn(3e7,1);
y = myfunc(x); % Copy
x = myfuncInPlace(x); % In-place
x = myfunc(x); % Copy
```

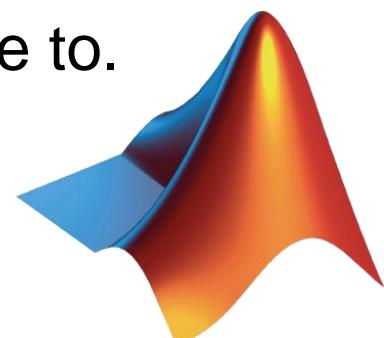


Agenda

- Leveraging the power of vector and matrix operations in MATLAB®
 - Demonstration: Preallocation and vectorization
- How does MATLAB® store and provide access to its' variables?
 - Demonstration: Row- vs. Column-major indexing
 - Demonstration: Logical indexing
- Addressing bottlenecks
 - Demonstration: Profiling code in MATLAB®
- Memory considerations
 - Demonstration: Copy-on-write
-  Summary and Q&A

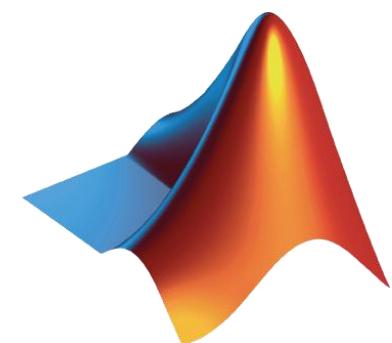
Summary

- Consider performance benefit of vector and matrix operations in MATLAB
- Preallocate memory and look at your indexing.
- Analyze your code for bottlenecks and address most critical items
- Never write hard-to-read code unless you have to.



Summary (continued...)

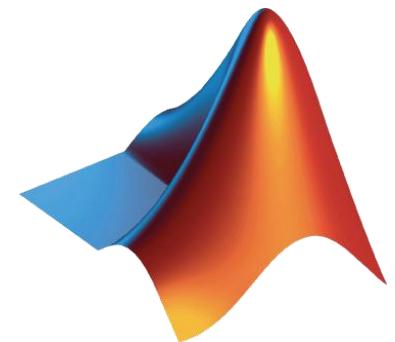
- MATLAB passes data to functions by value using lazy copy.



Sample of Other Performance Resources

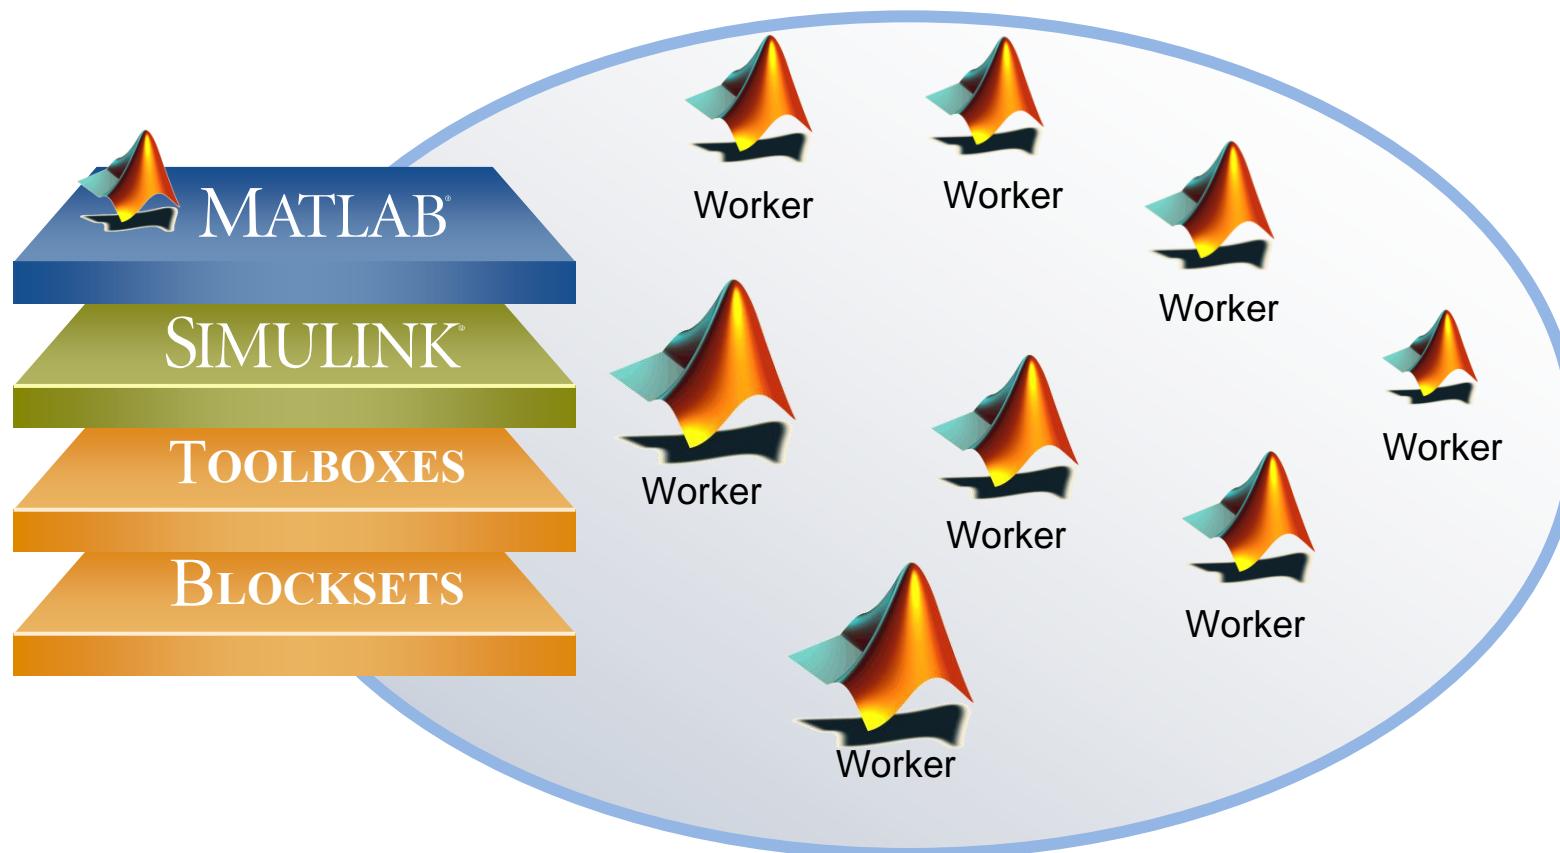
- MATLAB documentation
 - MATLAB → Programming Fundamentals → Performance
- Memory Management Guide
 - www.mathworks.com/support/tech-notes/1100/1106.html?BB=1
- The Art of MATLAB, Loren Shure's blog
 - blogs.mathworks.com/loren/
- MATLAB Answers
 - <http://www.mathworks.com/matlabcentral/answers/>

Parallel computing with MATLAB



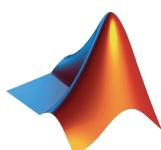
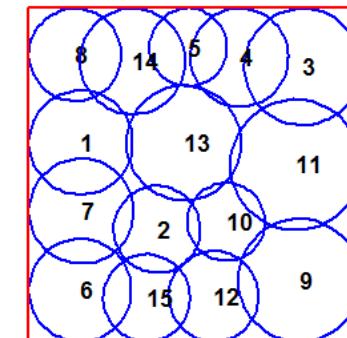
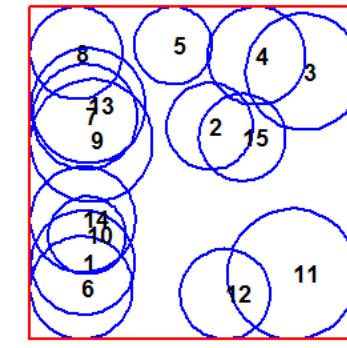
Going Beyond Serial MATLAB Applications

Use the power of multi-core machines and clusters



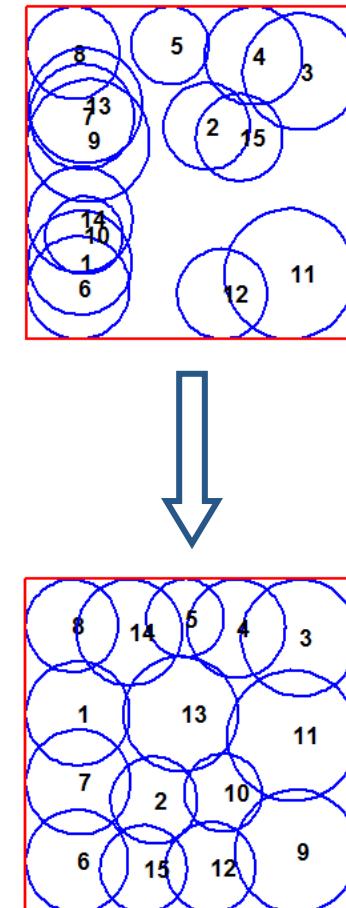
Example: Optimizing Tower Placement

- Determine location of cell towers
- Maximize coverage
- Minimize overlap



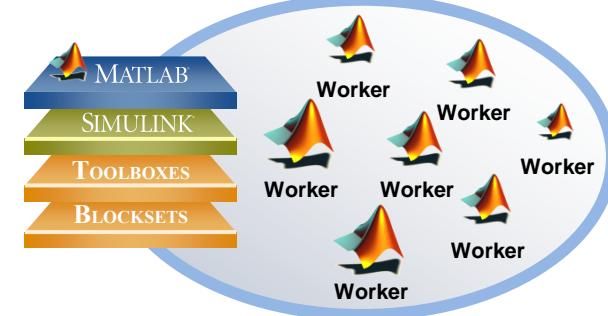
Summary of Example

- Enabled built-in support for Parallel Computing Toolbox in Optimization Toolbox
- Used a pool of MATLAB workers
- Optimized in parallel using `fmincon`



Tools Providing Parallel Computing Support

- Optimization Toolbox
- Global Optimization Toolbox
- Statistics Toolbox
- Communications System Toolbox
- Simulink Design Optimization
- Bioinformatics Toolbox
- Image Processing Toolbox
- ...



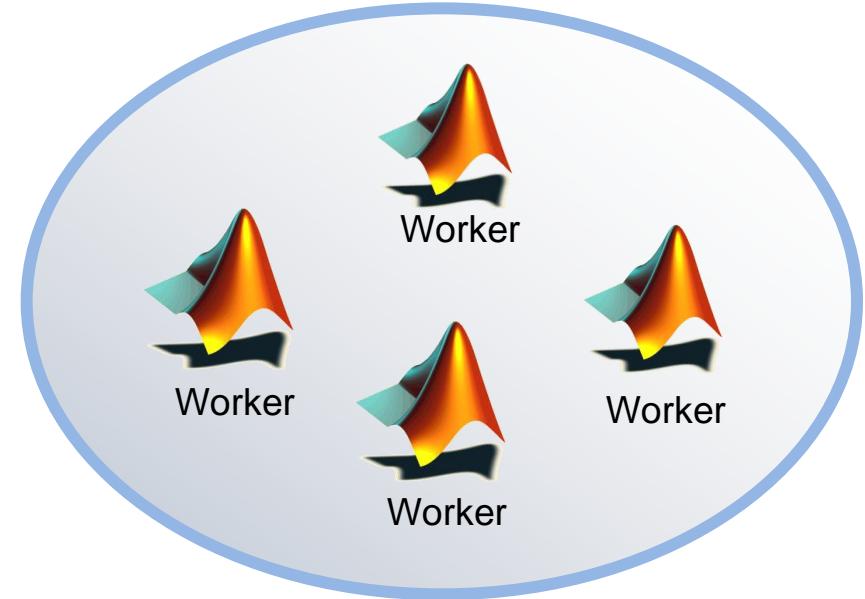
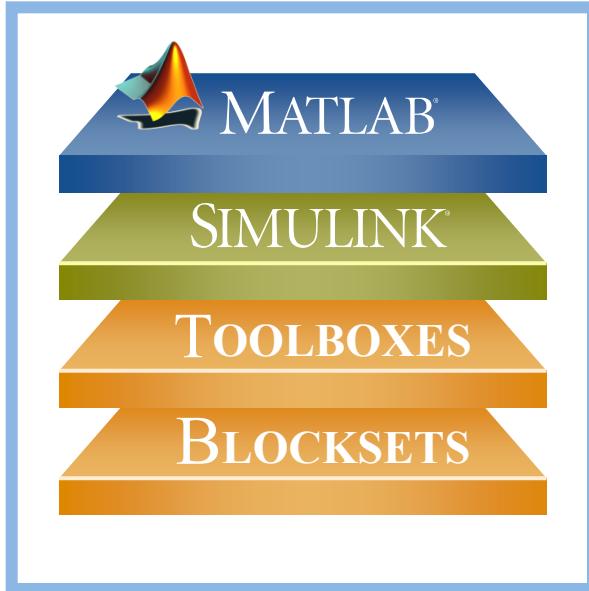
Directly leverage functions in Parallel Computing Toolbox

<http://www.mathworks.com/products/parallel-computing/builtin-parallel-support.html>

Agenda

-  Task parallel applications
 - GPU acceleration
 - Data parallel applications
 - Using clusters and grids

Task Parallel Applications

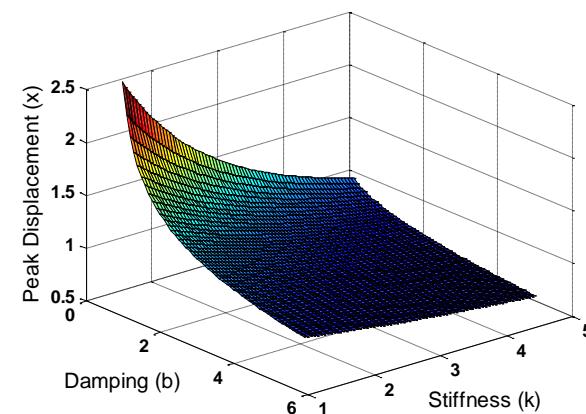
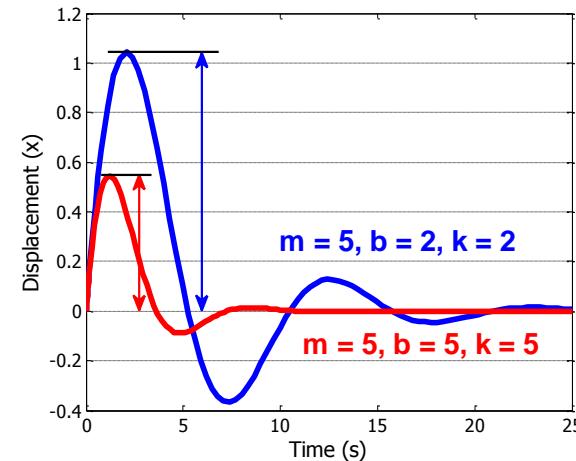
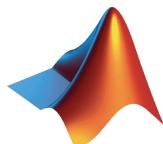


Example: Parameter Sweep of ODEs

- Solve a 2nd order ODE

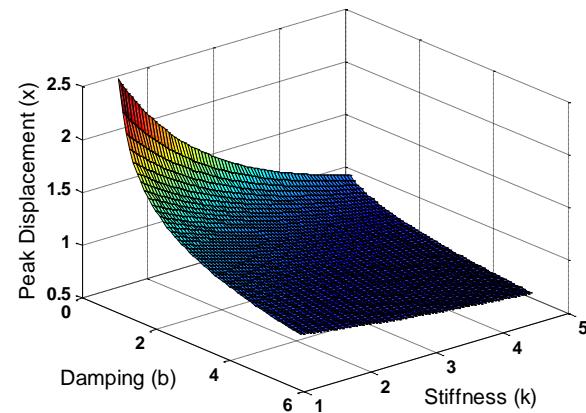
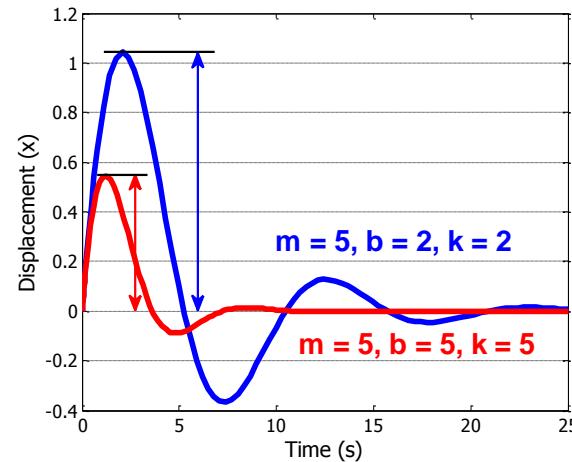
$$m\ddot{x} + \underbrace{b}_{1,2,\dots} \dot{x} + \underbrace{k}_{1,2,\dots} x = 0$$

- Simulate with different values for b and k
- Record peak value for each run
- Plot results

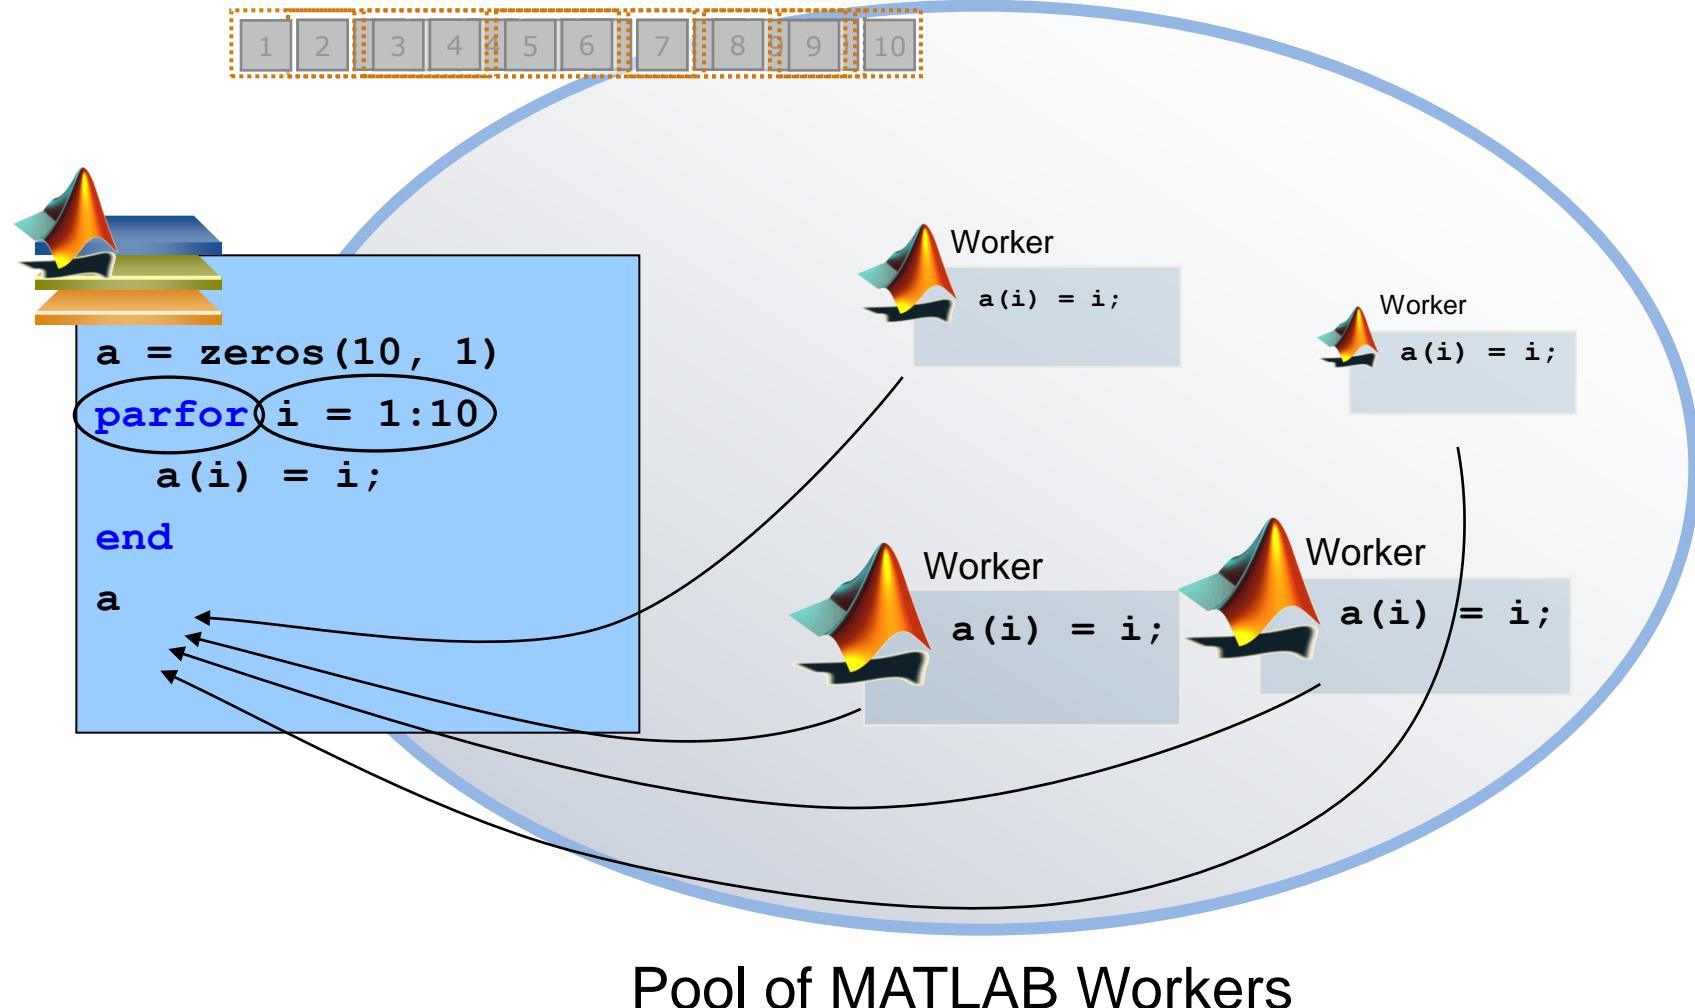


Summary of Example

- Mixed task-parallel and serial code in the same function
- Ran loops on a pool of MATLAB resources



The Mechanics of `parfor` Loops



Pool of MATLAB Workers

Converting `for` to `parfor`

- Requirements for `parfor` loops
 - Task independent
 - Order independent
- Constraints on the loop body
 - Cannot “introduce” variables (e.g. `load`, etc.)
 - Cannot contain `break` or `return` statements
 - Cannot contain another `parfor` loop

Advice for Converting `for` to `parfor`

- Use Code Analyzer to diagnose `parfor` issues
- If your `for` loop cannot be converted to a `parfor`, consider wrapping a subset of the body to a function
- Read the section in the documentation on classification of variables

Agenda

- Task parallel applications
- GPU acceleration
- Data parallel applications
- Using clusters and grids

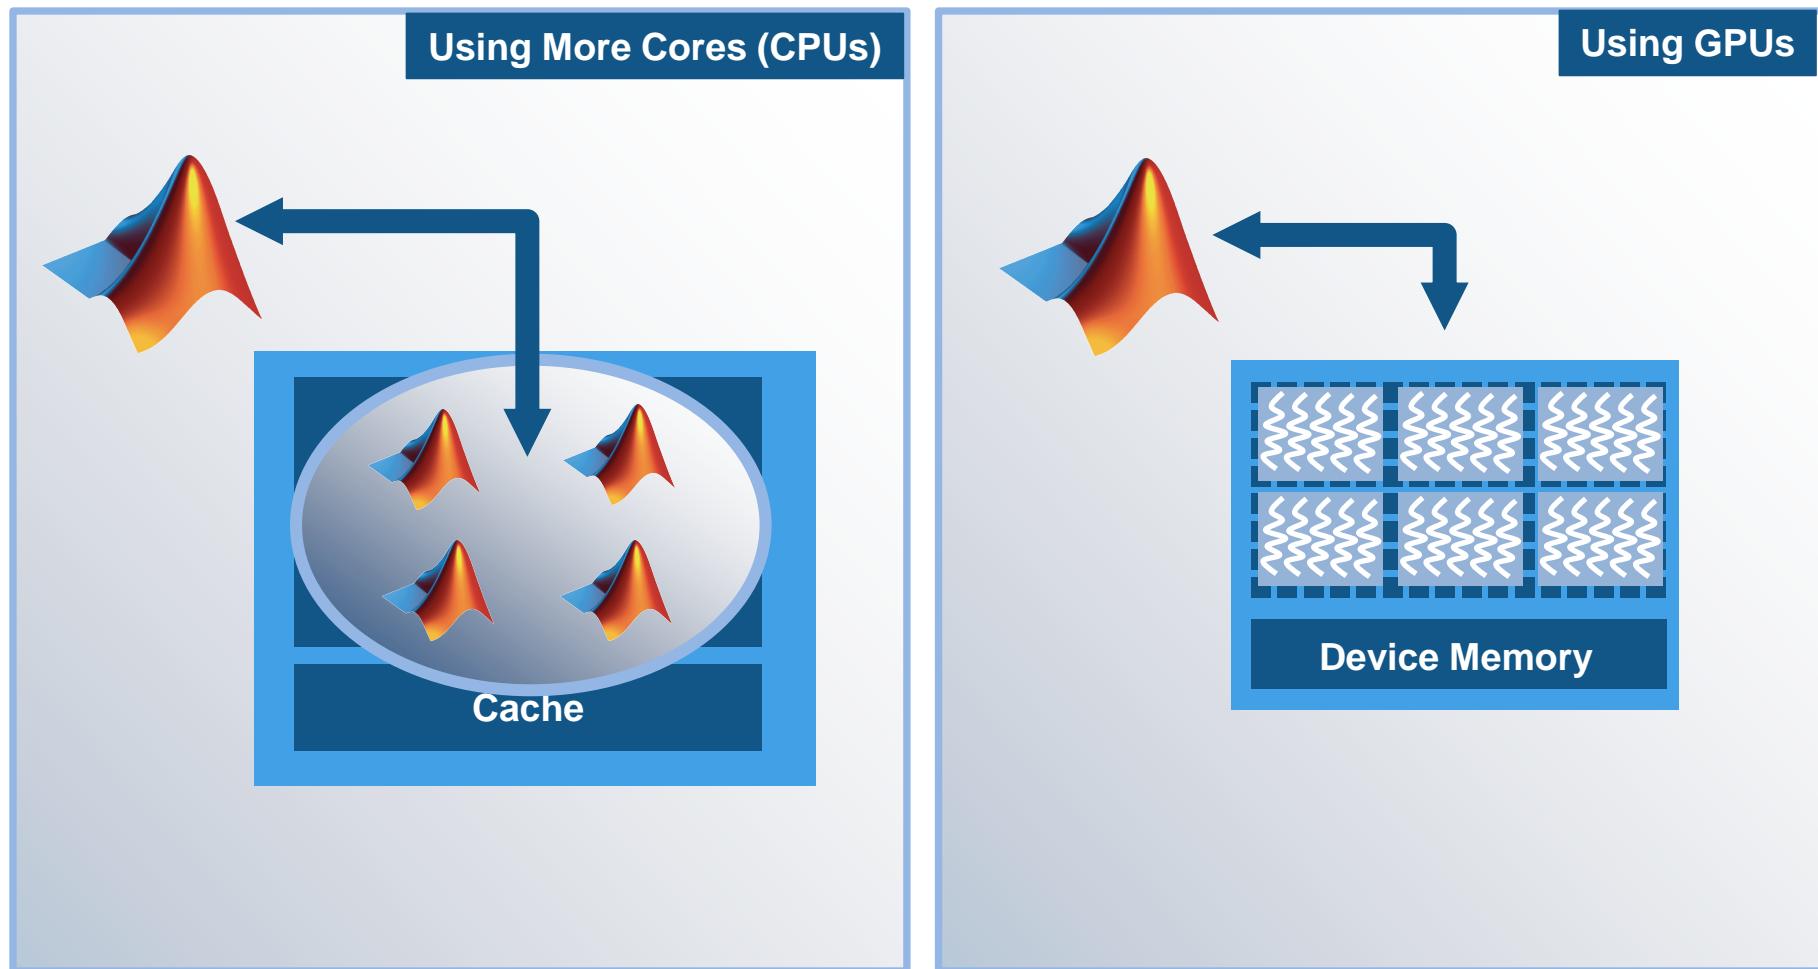


What is a Graphics Processing Unit (GPU)

- Originally for graphics acceleration, now also used for scientific calculations
 - Massively parallel array of integer and floating point processors
 - Typically hundreds of processors per card
 - GPU cores complement CPU cores
 - Dedicated high-speed memory
- * Parallel Computing Toolbox requires NVIDIA GPUs with Compute Capability 1.3 or greater, including NVIDIA Tesla 10-series and 20-series products. See http://www.nvidia.com/object/cuda_gpus.html for a complete listing

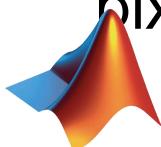
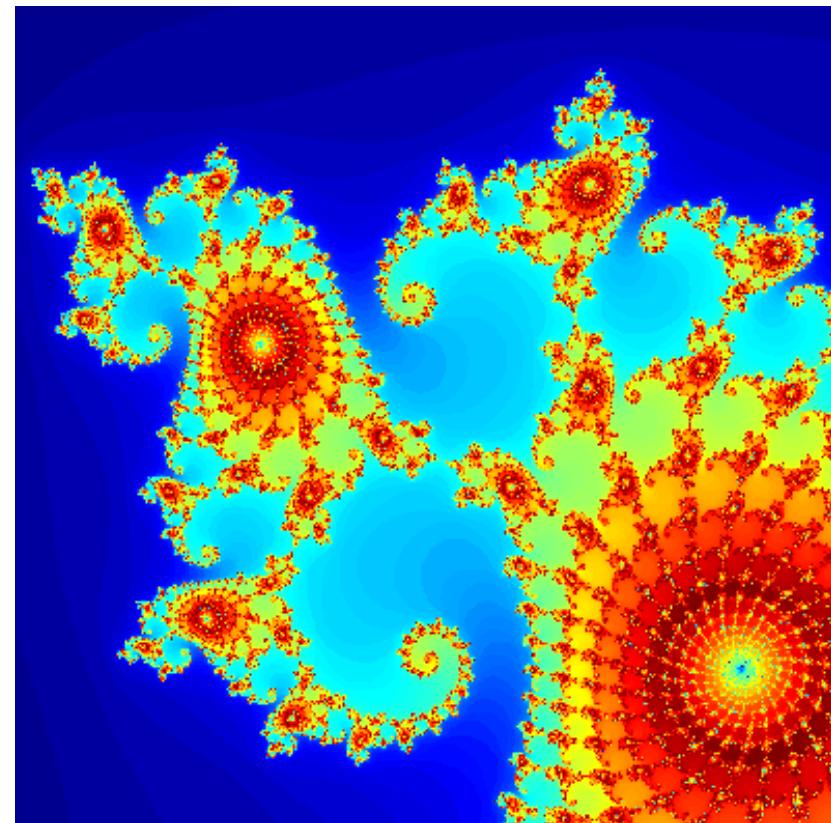


Performance Gain with More Hardware



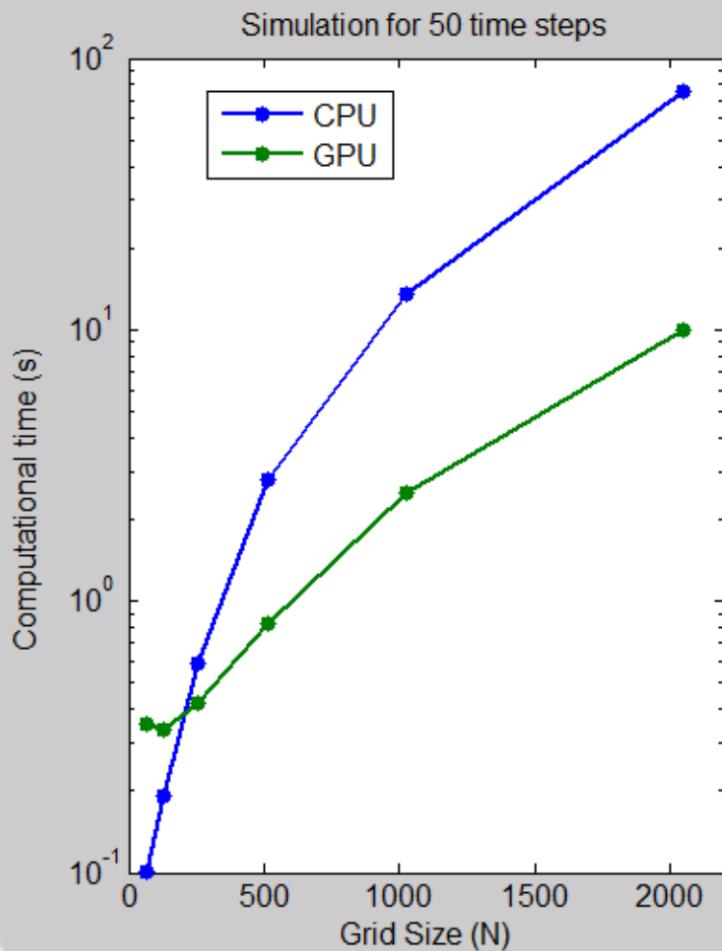
Example: Mandelbrot set

- The color of each pixel is the result of hundreds or thousands of iterations
- Each pixel is independent of the other pixels
- Hundreds of thousands of pixels



Real-world performance increase

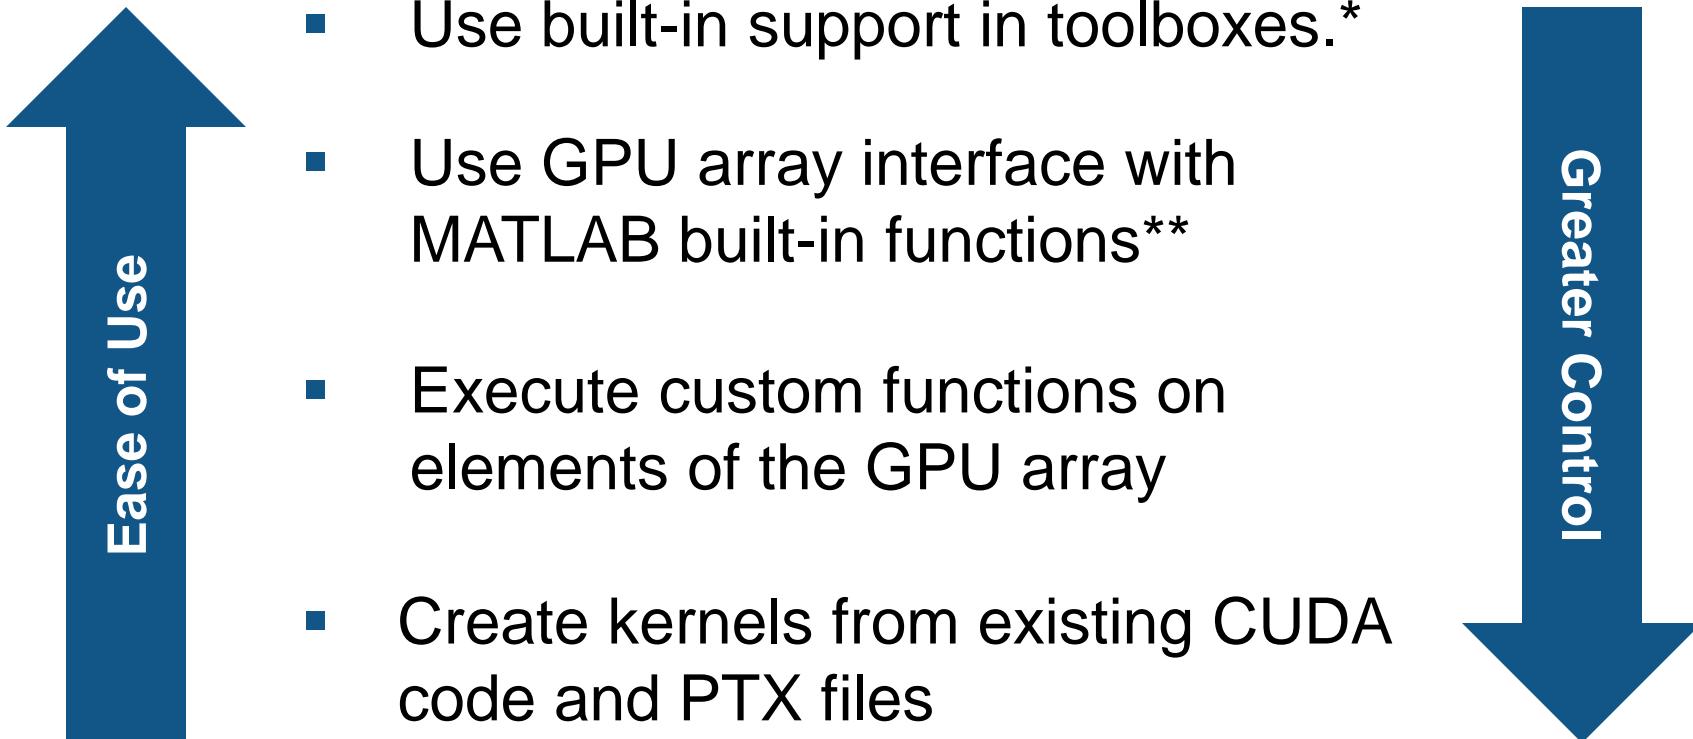
Solving a wave equation



Grid Size	CPU (s)	GPU (s)	Speedup
64 x 64	0.1004	0.3553	0.28
128 x 128	0.1931	0.3368	0.57
256 x 256	0.5888	0.4217	1.4
512 x 512	2.8163	0.8243	3.4
1024 x 1024	13.4797	2.4979	5.4
2048 x 2048	74.9904	9.9567	7.5

Intel Xeon Processor X5650, NVIDIA Tesla C2050 GPU

Summary of Options for Targeting GPUs



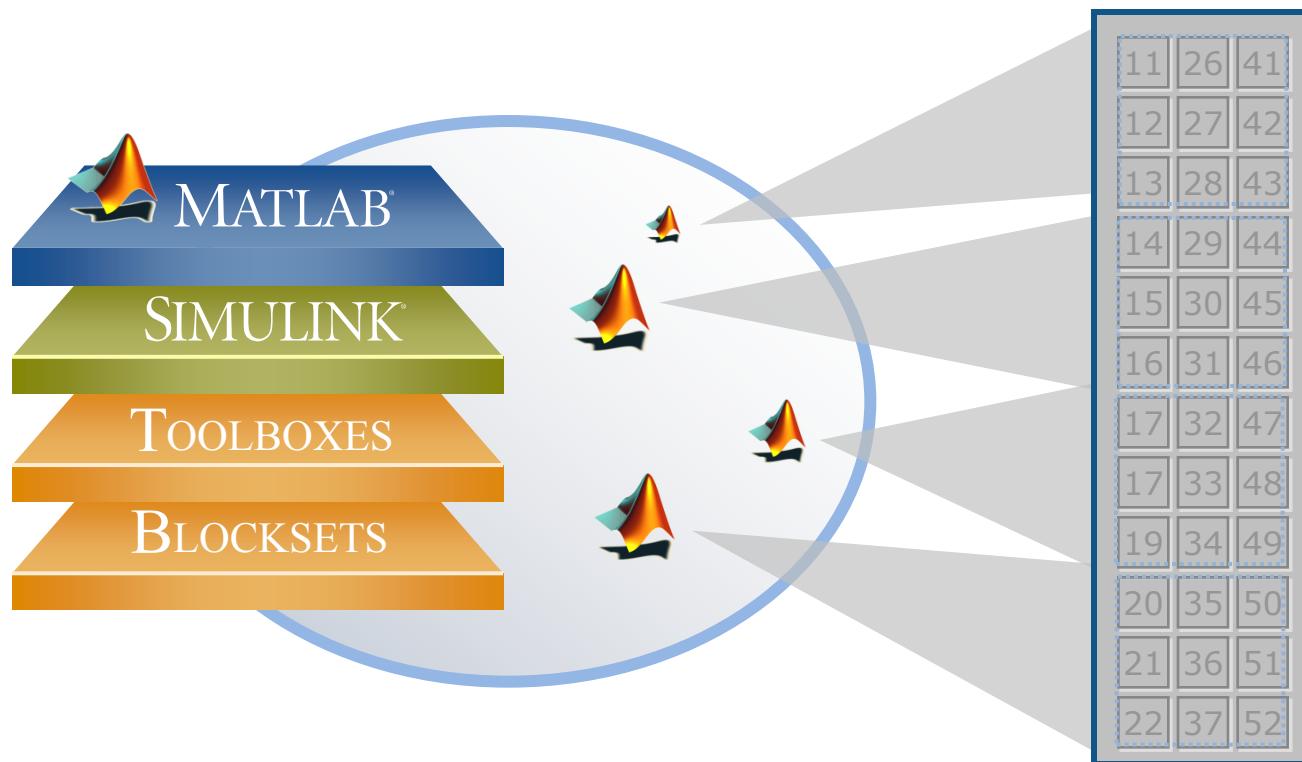
* Signal Processing Toolbox, Neural Network Toolbox and Phased Array Toolbox

** Over 200 functions support GPU Array data

Agenda

- Task parallel applications
- GPU acceleration
-  Data parallel applications
- Using clusters and grids

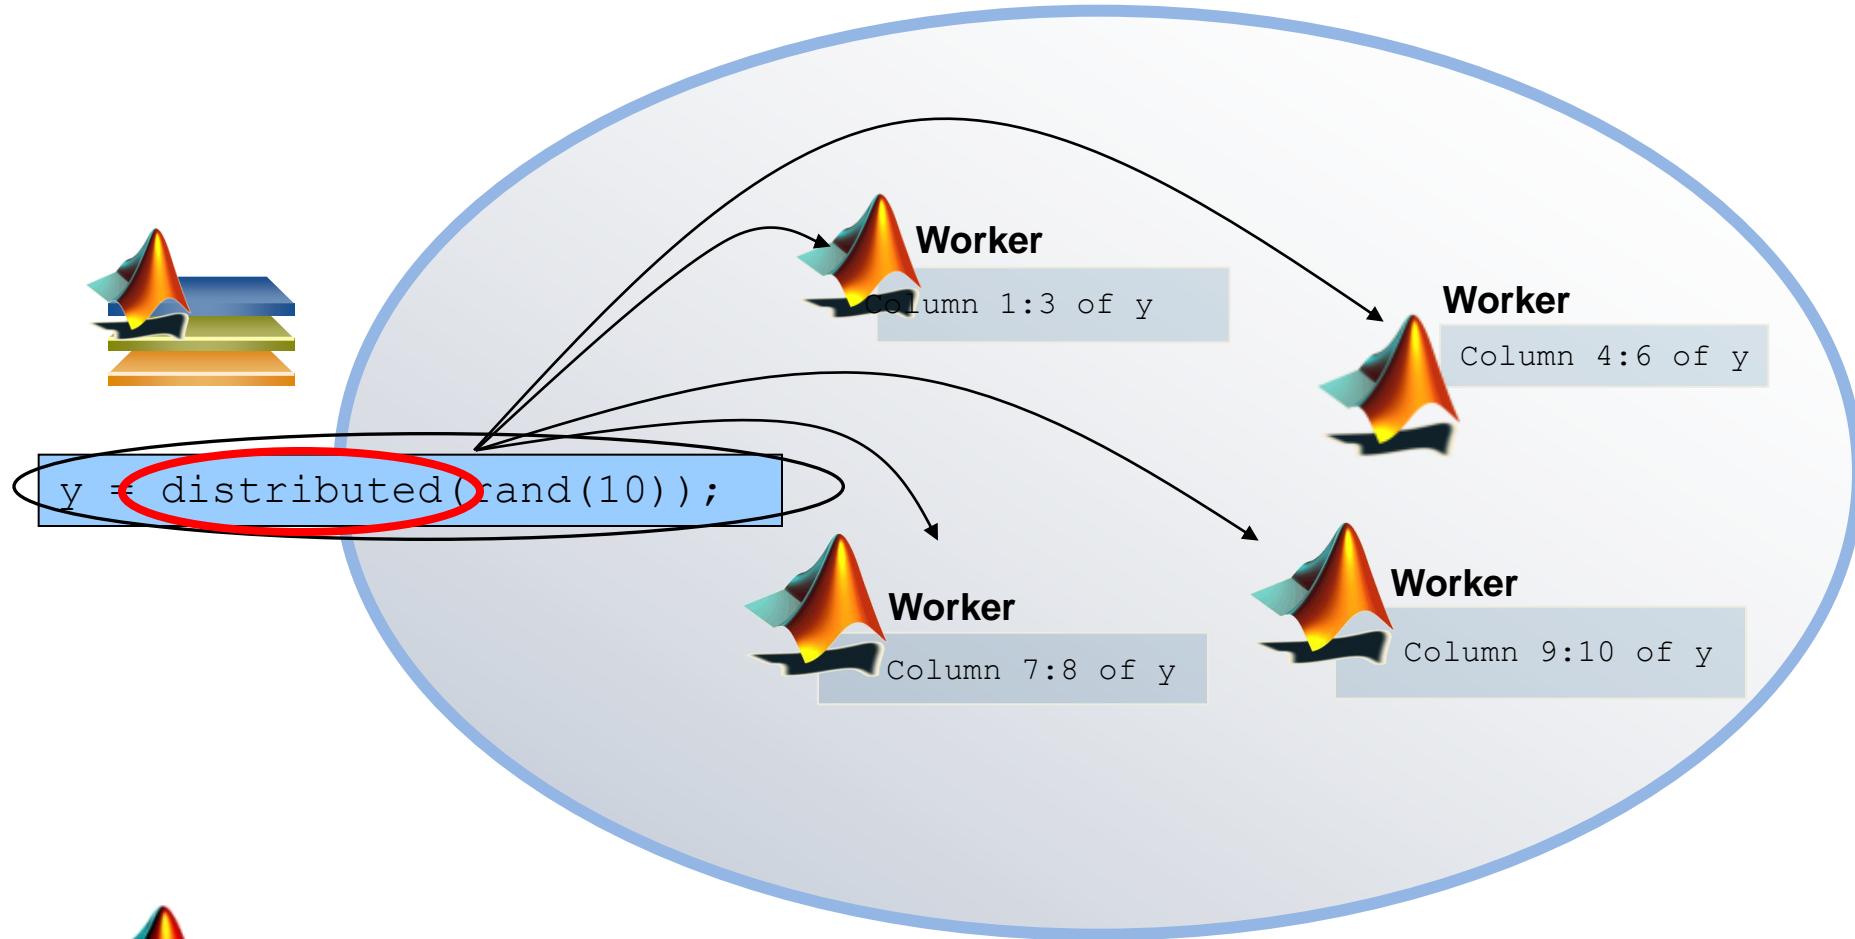
Big data: Distributed Arrays



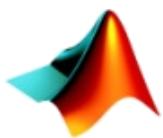
Remotely Manipulate Array
from Desktop

Distributed Array
Lives on the Cluster

Big Data: Distributed Arrays



Pool of MATLAB Workers

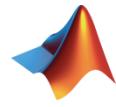
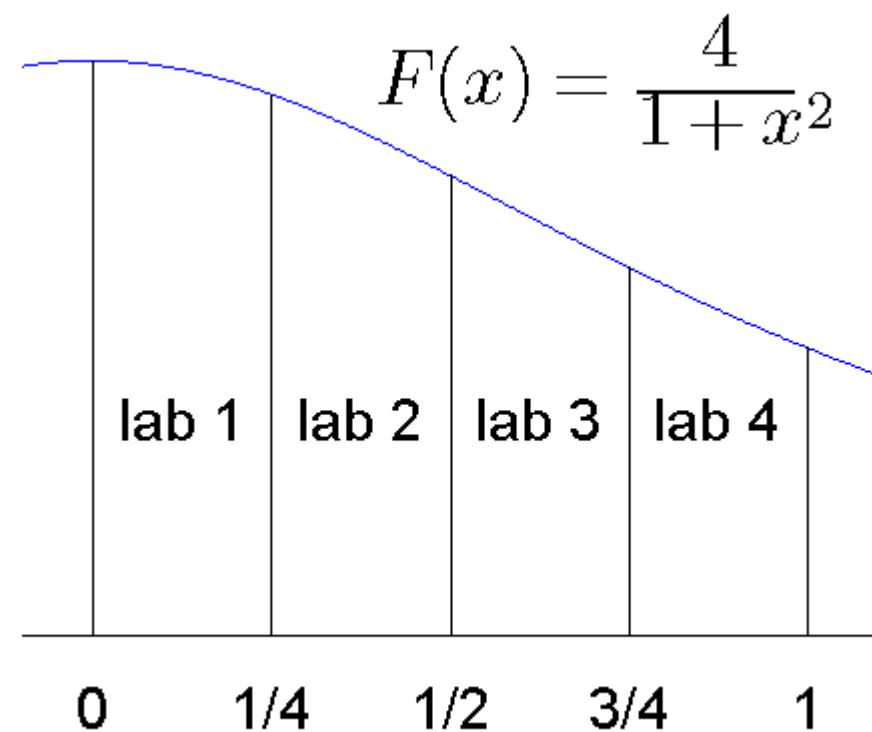


MATLAB Functions on Distributed Arrays ~ 150 supported functions

Type of Function	Function Names
Data functions	cumprod , cumsum , fft , max , min , prod , sum
Data type functions	arrayfun , cast , cell2mat , cell2struct , celldisp , cellfun , char , double , fieldnames , int16 , int32 , int64 , int8 , logical , num2cell , rmfield , single , struct2cell , swapbytes , typecast , uint16 , uint32 , uint64 , uint8
Elementary and trigonometric functions	abs , acos , acosd , acosh , acot , acotd , acoth , acsca , acsch , angle , asec , asecd , asech , asin , asind , asinh , atan , atan2 , atand , atanh , ceil , complex , conj , cos , cosd , cosh , cot , cotd , coth , csc , cscd , csch , exp , expml , fix , floor , hypot , imag , isreal , log , log10 , logip , log2 , mod , nextpow2 , nthroot , pow2 , real , reallog , realpow , realsqrt , rem , round , sec , secd , sech , sign , sin , sind , sinh , sqrt , tan , tand , tanh
Elementary matrices	cat , diag , eps , find , isempty , isequal , isequaln , isfinite , isinf , isnan , length , meshgrid , ndgrid , ndims , numel , reshape , size , sort , tril , triu
Matrix functions	chol , eig , inv , lu , norm , normest , qr , svd
Array operations	all , and ($\&$), any , bitand , bitor , bitxor , ctranspose ('), end , eq (==), ge (\geq), gt ($>$), horzcat ([]), ldivide (. \), le (\leq), lt ($<$), minus (-), mldivide (\), mrdivide (/), mtimes (*), ne (\neq), not (~), or (!), plus (+), power (. ^), rdivide (. /), subsasgn , subsindex , subsref , times (. *), transpose (. '), uminus (-), uplus (+), vertcat ([;]), xor
Sparse matrix functions	full , issparse , nnz , nonzeros , nzmax , sparse , spfun , spones
Special functions	dot

Demo: Approximation of π

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$



Interactive to Scheduling

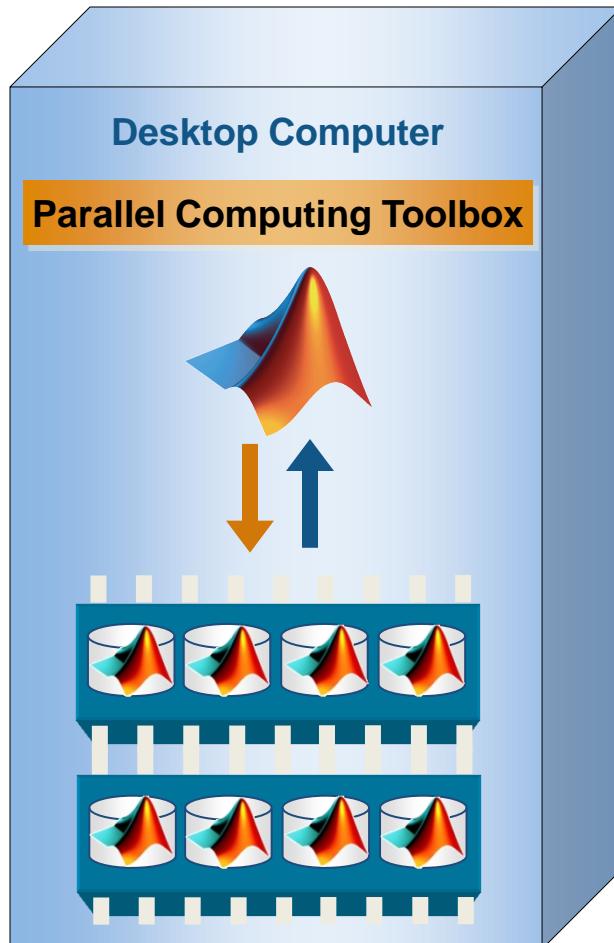
- Interactive
 - Great for prototyping
 - Immediate access to MATLAB workers
- Scheduling: **batch ('script.m') ;**
 - Offloads work to other MATLAB workers (local or on a cluster)
 - Access to more computing resources for improved performance
 - Frees up local MATLAB session

Agenda

- Task parallel applications
- GPU acceleration
- Data parallel applications
- Using clusters and grids



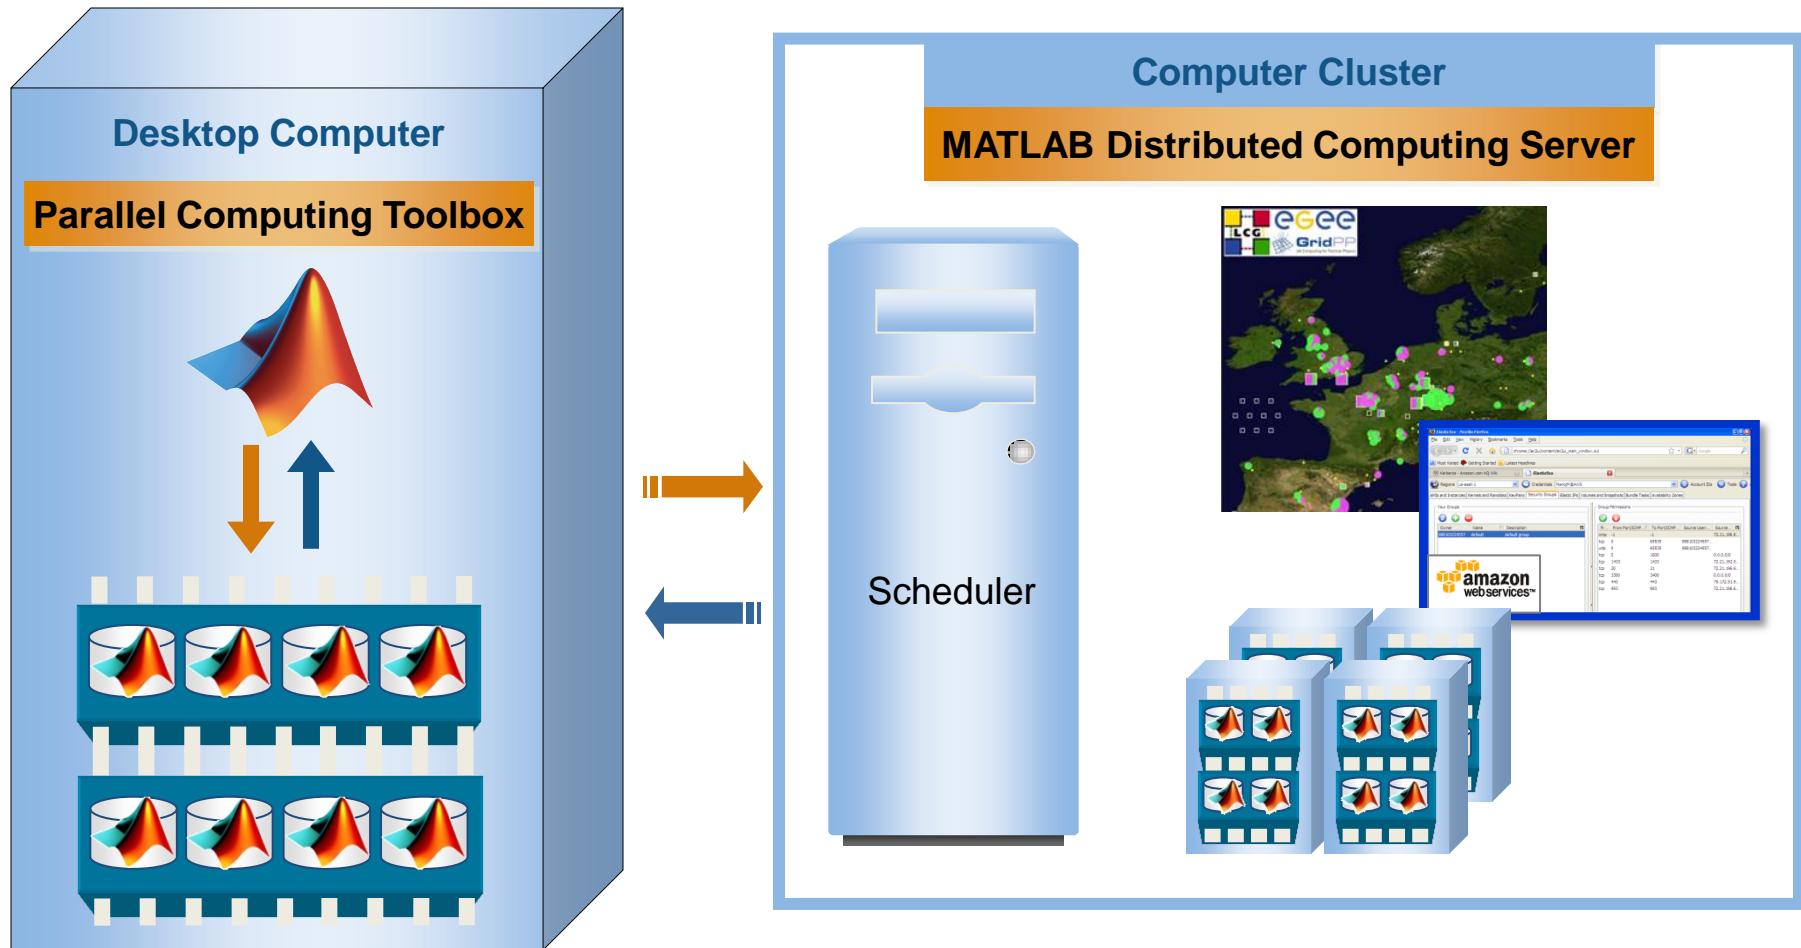
Run up to 12 Local Workers on Desktop



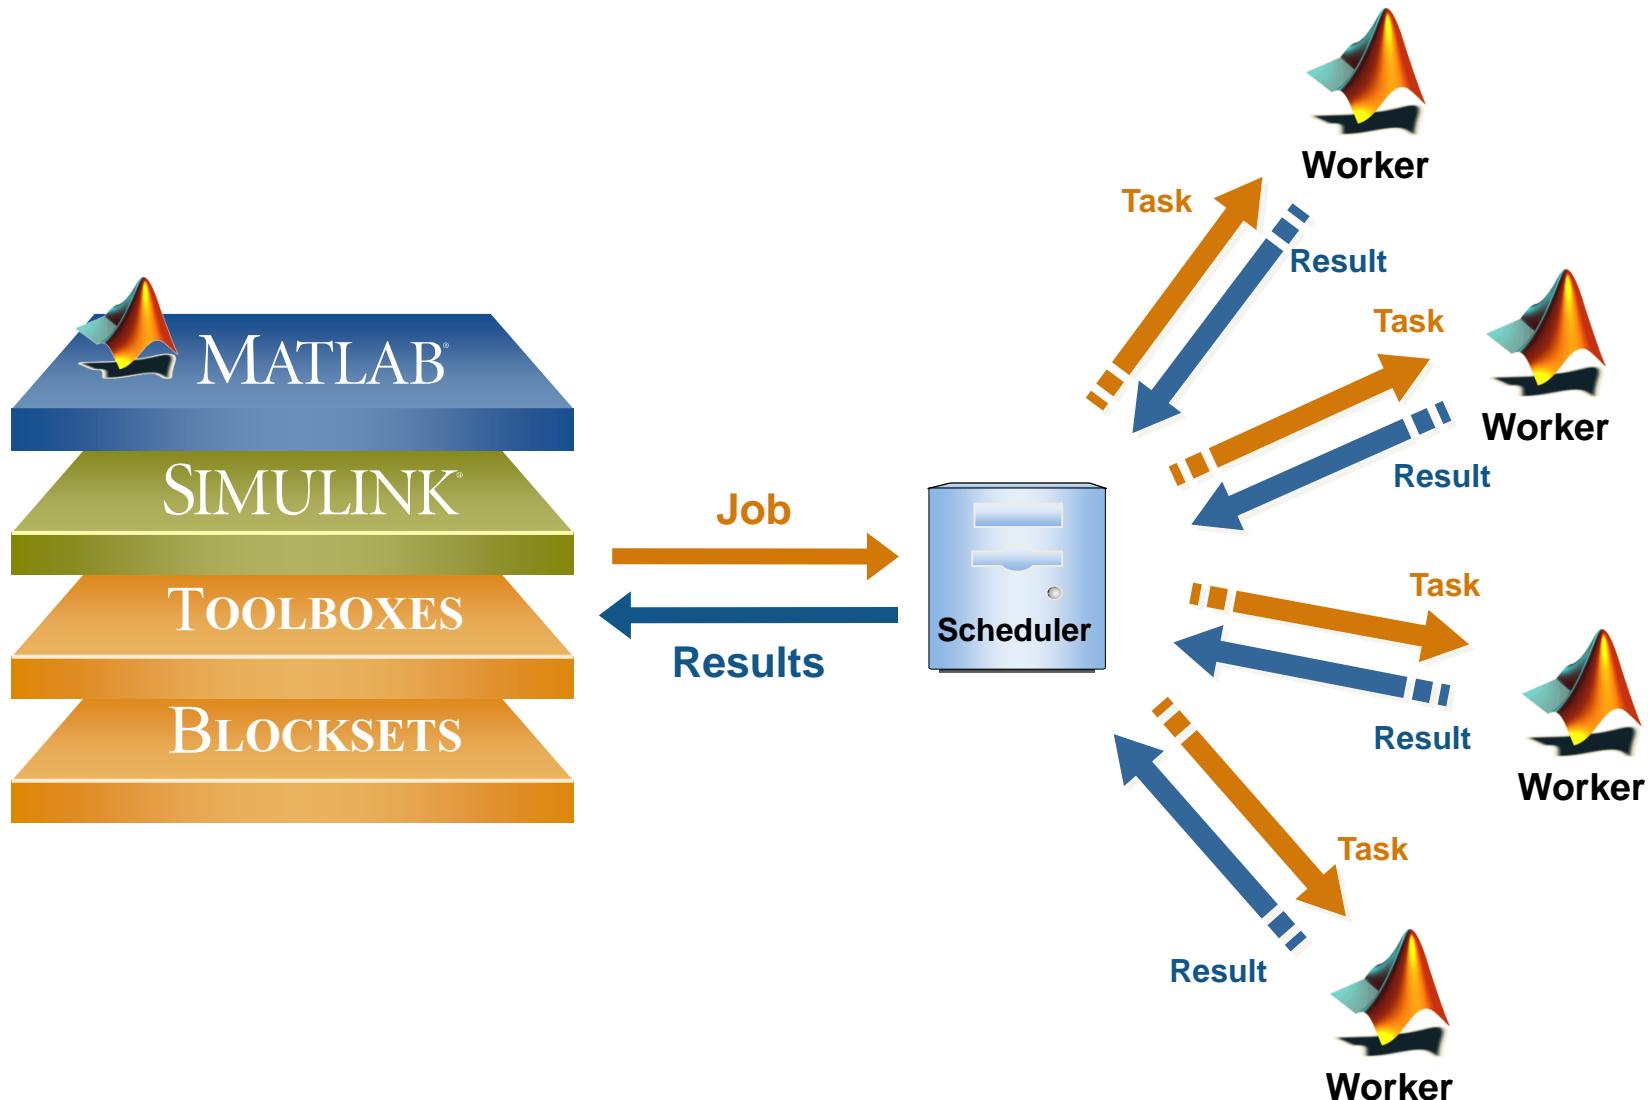
- Rapidly develop parallel applications on local computer
- Take full advantage of desktop power
- Separate computer cluster not required

Scale Up to Clusters, Grids and Clouds

No changes in your code



Scheduling Jobs and Tasks



Summary

- Developing parallel computations in MATLAB is easy.
- Parallel Computing Toolbox allows you to develop task-parallel, GPU, and data-parallel tasks
- Code developed using Parallel Computing Toolbox can run on clusters with minimal changes.
- GPU computing gives 5-10 times speedup with consumer hardware.

